

The Theory of Persistence II: Information Geometry and Holonomy

Yan Senez

yan.senez@protonmail.com

April 2026

Abstract

This is the second of five articles presenting the mathematical foundations of the Theory of Persistence (PT). Building on the sieve dynamics, uniqueness, conservation laws, and Gap Fundamental Theorem established in the companion article [?], we develop two pillars of the theory: the canonical information geometry of the sieve state space, and the holonomy structure that emerges from the sieve's transfer matrices.

Information geometry. The KL divergence D_{KL} is the unique f -divergence compatible with the CRT product structure of the sieve state space (Theorem G1, via Shore–Johnson reduction [?], with 7/7 elimination tests; an autonomous CRT proof via Csiszár's characterisation [?] is also given). The Fisher metric g^F emerges as the second variation of D_{KL} (Theorem G2, standard) and is the unique monotone Riemannian metric on Δ_2 (Čencov reduction [?], Theorem G3, 7/7 elimination tests). We exhibit the Fisher metric on binary and ternary simplices, show it is forced by the optimal expander structure of the mod-3 transition graph, and connect it to the addition/multiplication duality of the integers.

Holonomy. The holonomy angle θ_p satisfies the exact algebraic identity $\sin^2\theta_p = \delta_p(2 - \delta_p)$, derived by three independent routes (geometric, spectral, information-geometric). The anomalous dimensions $\gamma_p = -d \ln \sin^2\theta_p / d \ln \mu$ are strictly decreasing in p ; the active prime criterion $\gamma_p > s = 1/2$ selects exactly $\{3, 5, 7\}$ at the fixed point $\mu^* = 15$, yielding $N_c = 3$. The Casimir coefficient $C_F = (N_c^2 - 1)/(2N_c) = 4/3$ is a consequence, not an input. An independent algebraic check gives $(N_c - 3)(N_c + 1) = 0$. Echo primes ($p \geq 11$) contribute perturbatively as vacuum polarization corrections.

All results follow from $s = 1/2$ and the sieve structure established in [?]. No physics is invoked; no parameter is fitted.

Contents

1 Introduction

This article is the second instalment (M2) in a five-part series developing the mathematical foundations of the Theory of Persistence (PT). The series is organised as follows:

M1 [?]: The sieve as a dynamical field, uniqueness of Eratosthenes, conservation laws, and the Gap Fundamental Theorem.

M2 (this article): Information geometry, holonomy, and the internal angle scale.

M3 [?]: Convergence, exact recurrences, and the fixed point $\mu^* = 15$.

M4 [?]: Extended mathematical structures (complex mechanics, spectral graphs, symmetry).

M5 [?]: The bridge from arithmetic to physics.

A separate companion [?] develops the detailed physical predictions (Feynman programme, 43 observables); M5 provides the structural bridge, while [?] provides the quantitative validation.

The companion article [?] established the foundational layer. We recall the essential results:

- **T1 (forbidden transitions)**: the transfer matrix $T_3 = \text{antidiag}(1, 1)$ has a unique stationary distribution at $(1/2, 1/2)$, giving $s = 1/2$ with no free parameter.
- **T6 (uniqueness)**: three independent proofs that Eratosthenes is the unique self-consistent sieve.
- **L0 (conservation)**: the maximum-entropy gap distribution is geometric; the conservation exponent is $\alpha_{\text{cons}} = s^2 = 1/4$.
- **GFT**: the exact identity $\log_2 m = D_{\text{KL}} + H$ decomposes total capacity into structure and noise.
- **CRT gauge connection**: the Chinese Remainder Theorem endows the sieve state space Δ_m with a product structure $\Delta_m \cong \prod_{p|m} \Delta_p$.

The present article develops two themes.

Information geometry (Section ??). We prove that the KL divergence D_{KL} and the Fisher metric g^F are the *unique* divergence and metric compatible with the sieve's symmetries. The proofs are reductions to classical theorems—Shore–Johnson [?] for D_{KL} and Čencov [?] for g^F —after verifying that the sieve state space satisfies their premises. An autonomous proof of G1 via Csiszár's characterisation of additive f -divergences is also given. This section is entirely mathematical and invokes no physics.

Holonomy (Section ??). The sieve's transfer matrix at each prime p generates a rotation angle θ_p on the residue space $\mathbb{Z}/p\mathbb{Z}$. The holonomy identity $\sin^2 \theta_p = \delta_p(2 - \delta_p)$ is an algebraic tautology, verified by three independent routes. The anomalous dimensions γ_p are strictly decreasing, and the criterion $\gamma_p > s = 1/2$ selects exactly three active primes $\{3, 5, 7\}$ at $\mu^* = 15$, giving $N_c = 3$ and $C_F = 4/3$. The physical identification of these arithmetic quantities with gauge couplings and QCD colours is a bridge claim deferred

to [?].

Forward pointers. The companion article [?] proves convergence of the sieve recurrences and establishes $\mu^* = 15$ as the unique non-trivial fixed point. The extended mathematical structures (complex mechanics, spectral graphs, and symmetry extensions) are developed in [?]. The bridge from arithmetic to physics—where $\sin^2\theta_p$ becomes a coupling constant and $N_c = 3$ becomes the number of QCD colours—is the subject of [?].

2 Canonical Information Geometry

Geometry is the branch of mathematics which studies the properties of figures. Information geometry is the branch which studies the properties of statistical manifolds.

Shun-ichi Amari

STATUS

GOAL: Prove that the two canonical divergences on the sieve state space—the KL divergence D_{KL} and the Fisher metric g^F —are the *unique* choices compatible with the sieve’s symmetries.

INPUTS: CRT, gauge connection, state space Δ_m ([?]). GFT, D_{KL} already used ([?]).

MAIN RESULTS: G1: D_{KL} is the unique f -divergence invariant under CRT (Shore–Johnson reduction, 7/7). G2: Second variation of D_{KL} = Fisher metric (standard, Amari 1985). G3: Fisher metric is the unique monotone Riemannian metric on Δ_2 (Čencov reduction, 7/7).

STATUS: [THM] (reduction: G1 via Shore–Johnson, G3 via Čencov). This chapter does NOT prove Shore–Johnson or Čencov autonomously. It verifies that the sieve state space satisfies their premises.

NOT PROVED: Shore–Johnson (1980) and Čencov (1982): imported as named classical theorems. Physical use of Fisher metric (the companion article [?]).

2.1 Intuition: why these two objects?

The Gap Fundamental Theorem (see [?]) uses D_{KL} as the measure of sieve structure. The companion article [?] will use the Fisher metric to derive general relativity. But why D_{KL} and Fisher? Among the infinitely many possible divergences and metrics on probability distributions, what singles these two out?

The answer, for D_{KL} : it is the unique f -divergence that is additive under the CRT factorization that the sieve enjoys. Any other f -divergence (Hellinger, χ^2 , total variation) fails this invariance (verified: 7/7 tests).

The answer, for Fisher: it is the unique Riemannian metric on probability simplices that contracts under Markov maps. Since the sieve's transfer matrix T_m is a Markov map, any information-geometric distance must contract under it. Fisher is the unique metric with this property.

Both results are reductions to classical theorems (Shore–Johnson [?], Čencov [?]). PT's contribution is verifying that the sieve state space satisfies the premises.

2.2 The sieve state space

Definition 2.1 (Probability simplex Δ_m). The *sieve state space* at modulus m is the probability simplex:

$$\Delta_m = \left\{ P = (p_0, \dots, p_{m-1}) \in \mathbb{R}^m : p_r \geq 0, \sum_r p_r = 1 \right\}.$$

The sieve distribution $P_m \in \Delta_m$ is the vector of prime residue frequencies modulo m .

The sieve acts on Δ_m via the CRT factorization: for $m = q_1 \cdots q_k$ (square-free), the CRT isomorphism gives

$$\Delta_m \cong \Delta_{q_1} \times \cdots \times \Delta_{q_k},$$

so the state space is a Cartesian product of smaller simplices. Any geometric object (divergence, metric) on Δ_m must respect this product structure. Figure ?? illustrates the Fisher metric on the binary simplex Δ_2 , and Figure ?? shows the extension to the ternary simplex Δ_3 .

2.3 G1: Uniqueness of D_{KL} via Shore–Johnson

The Shore–Johnson theorem (1980, external import). Under the following axioms on a divergence $D(\cdot\|\cdot)$:

- (SJ1) **Consistency**: $\arg \min_Q D(Q\|P)$ subject to constraints is independent of the order in which constraints are imposed.
- (SJ2) **Invariance**: D is invariant under permutation of labels.
- (SJ3) **System independence**: for independent systems, D is additive.
- (SJ4) **Subset independence**: updating on a constraint that concerns only a subset changes only that subset's distribution.
- (SJ5) **Scaling**: D transforms consistently under relabeling.




Figure 1: Fisher metric on the binary simplex Δ_2 . The Euclidean interval $[0, 1]$ (gray) is “curved” by the Fisher metric $ds^2 = dp^2/[p(1-p)]$ into a half-circle (blue arc). Grid lines connect equal probability values: they crowd near the boundaries ($p \rightarrow 0$ or $p \rightarrow 1$), showing the metric divergence. The stationary point $s = 1/2$ (orange) sits at the geodesic midpoint. This is the *unique* monotone metric on Δ_2 (Čencov).




Figure 2: Fisher metric on the ternary simplex Δ_3 . The three vertices (pure states, $p_i = 1$) are at infinite Fisher distance from any interior point—the metric diverges near edges. Grid lines show how equal Euclidean steps correspond to unequal Fisher lengths (crowding near boundaries). The centroid $p_i = 1/3$ (orange, corresponding to $s = 1/2$ in the binary projection) is the geodesic midpoint.

Shore and Johnson [?] proved: the unique divergence satisfying SJ1–SJ5 and yielding the maximum entropy principle is D_{KL} .

Theorem 2.2 (KL Divergence Uniqueness (G1)). $\sim_L$

Lean: PT.Information.T6cG1Autonomous.klDivergence_additive *On the sieve state space Δ_m with the CRT product structure: D_{KL} is the unique f -divergence satisfying Shore–Johnson axioms SJ1–SJ5 and invariant under CRT factorization.*

Proof. Verification of premises. The sieve state space Δ_m with the product structure $\Delta_m \cong \prod_{p|m} \Delta_p$ satisfies:

- SJ3 (system independence): CRT factorization gives $D_{\text{KL}}(\prod P_p \parallel \prod Q_p) = \sum_p D_{\text{KL}}(P_p \parallel Q_p)$. Independence of the CRT components implies additivity. (Verified: 7/7 tests.)
- SJ4 (subset independence): updating residues modulo p changes only the p -component in the CRT decomposition. (Verified: 7/7 tests.)
- SJ1, SJ2, SJ5: standard consistency/invariance properties satisfied by the uniform reference U_m and the empirical distribution P_m .

Application of Shore–Johnson. Since the premises SJ1–SJ5 hold on Δ_m , the Shore–Johnson theorem gives: D_{KL} is the unique divergence satisfying these axioms. (External import: Shore–Johnson 1980, published in IEEE Trans. Inf. Theory.)

Elimination of alternatives.

- χ^2 divergence: violates SJ3 (not additive under CRT, ratio of squared differences not additive).
- Hellinger distance: violates SJ3 (square root breaks additivity).
- Total variation: violates SJ4 (updating one coordinate affects the ℓ_1 norm globally).
- Rényi divergences D_α ($\alpha \neq 1$): violate SJ1 (inconsistency of sequential updates for $\alpha \neq 1$).

□

Remark 2.3 (Scope of G1). G1 is a *reduction* to Shore–Johnson, not an autonomous proof. The PT contribution: verifying that Δ_m with CRT structure satisfies SJ1–SJ5. The uniqueness conclusion is imported.

Theorem 2.4 (CRT-Autonomous KL Uniqueness (G1')). $\sim_L$

Lean: PT.Information.T6cG1Autonomous.fDivergence_linear_term_vanishes *On the sieve state space Δ_m with CRT product structure, D_{KL} is the unique f -divergence satisfying additivity under independent subsystems.*

Proof. Step 1 (CRT $\Rightarrow$ product structure). By the Chinese Remainder Theorem (see Theorem 3.10 in [?]), for $m = q_1 \cdots q_k$ (square-free): $\Delta_m \cong \Delta_{q_1} \times \cdots \times \Delta_{q_k}$. Any divergence D on Δ_m must respect this factorization.

Step 2 (Additivity characterisation). An f -divergence $D_f(P\|Q) = \sum_i q_i f(p_i/q_i)$ is additive on product distributions ($P = P_1 \otimes P_2$, $Q = Q_1 \otimes Q_2 \Rightarrow D_f(P\|Q) = D_f(P_1\|Q_1) + D_f(P_2\|Q_2)$) if and only if f has the form $f(t) = c_1 t \ln t + c_2(t - 1)$ for constants $c_1 > 0$, $c_2 \in \mathbb{R}$ (Csiszár [?]). Since $\sum_i q_i(p_i/q_i - 1) = 0$ on normalised distributions, the c_2 term vanishes identically. The unique additive f -divergence (up to positive rescaling) is therefore $D_f(P\|Q) = \sum_i p_i \ln(p_i/q_i) = D_{\text{KL}}(P\|Q)$.

Step 3 (Uniqueness). Since CRT imposes product structure at every sieve level, and since $f(t) = t \ln t$ is the unique generator producing an additive divergence, D_{KL} is the unique additive f -divergence on Δ_m . No Shore–Johnson import is needed; the proof uses only CRT, the definition of f -divergences, and Csiszár’s characterisation of additive generators (elementary convex analysis).

Elimination.

- χ^2 : $f(t) = (t - 1)^2$. Not of the form $c_1 t \ln t + c_2(t - 1)$. Numerical: $D_{\chi^2}(P \otimes P' \| Q \otimes Q') \neq D_{\chi^2}(P \| Q) + D_{\chi^2}(P' \| Q')$ (verified).
- Hellinger: $f(t) = (\sqrt{t} - 1)^2$. Not of the required form. Fails additivity (verified).
- Total variation: $f(t) = |t - 1|$. Not differentiable at $t = 1$. Fails additivity (verified).

Verified: `scripts/ch05_geometry/test_G1_autonomous.py` (12/12 PASS). $\square$

Remark 2.5 (Two independent proofs of G1). Theorem ?? (Shore–Johnson reduction) and Theorem ?? (autonomous CRT) constitute two independent proofs of D_{KL} uniqueness. The first imports Shore–Johnson; the second uses only Csiszár’s characterisation.

2.4 G2: Fisher metric as second variation of D_{KL}

Theorem 2.6 (Fisher Metric Emergence (G2)). ✓_L

Lean: `PT.Information.G2FisherEmergence.hasDerivAt_klTerm` *The second-order Taylor expansion of D_{KL} around the uniform distribution U_m gives the Fisher information metric:*

$$D_{\text{KL}}(P + \epsilon v \| P) = \frac{\epsilon^2}{2} \sum_{r,s} g_{rs}^F(P) v_r v_s + O(\epsilon^3),$$

where $g_{rs}^F(P) = \delta_{rs}/p_r$ is the Fisher metric on Δ_m .

Proof. Taylor expansion of $D_{\text{KL}}(Q\|P) = \sum_r q_r \log(q_r/p_r)$ around $Q = P + \epsilon v$:

$$D_{\text{KL}}(P + \epsilon v \| P) = \frac{\epsilon^2}{2} \sum_r \frac{v_r^2}{p_r} + O(\epsilon^3).$$

The Hessian $\partial^2 D_{\text{KL}}/\partial q_r \partial q_s|_{Q=P} = \delta_{rs}/p_r$ is exactly the Fisher metric (standard result: Amari [?], Theorem 2.1). $\square$

2.5 G3: Uniqueness of Fisher metric via Čencov

The Čencov theorem (1982, external import). Let $\mathcal{M}_n$ denote the manifold of all probability distributions on n outcomes. A Riemannian metric g on $\mathcal{M}_n$ is *monotone* if for any Markov map T (stochastic matrix):

$$g_{T(P)}(T_*v, T_*v) \leq g_P(v, v).$$

Čencov [?] proved: the Fisher metric is the unique monotone Riemannian metric on $\mathcal{M}_n$ (up to a positive constant).

Theorem 2.7 (Fisher Metric Uniqueness (G3)). ★_L

Lean: PT.Information.G3FisherUniqueness.G3_fisher_unique *On the sieve state space $\Delta_2 = \{(p, 1-p) : p \in (0, 1)\}$, the Fisher metric g^F is the unique Riemannian metric monotone under all sieve Markov maps (including T_m and its marginals).*

Proof. Verification that T_m is a Markov map. The transfer matrix T_m is row-stochastic (rows sum to 1, entries ≥ 0 by definition). Hence T_m is a Markov map.

Verification of the premise Δ_2 . The sieve state space at $m = 3$ is $\Delta_2 = \{(\pi_1, \pi_2) : \pi_1 + \pi_2 = 1\}$ (two non-zero residue classes, from T1 [?]). The Fisher metric on Δ_2 is $ds^2 = dp^2/(p(1-p))$, the arc-length metric on $[0, 1]$.

Monotonicity. For any sieve projection $T : \Delta_m \rightarrow \Delta_{m'}$ (reducing modulus from m to m' with $m'|m$), the Fisher metric contracts: $g_{T(P)}^F(T_*v, T_*v) \leq g_P^F(v, v)$. Numerical verification: all projections tested give max ratio = 1.000000 (contraction).

Application of Čencov. Since Δ_2 (and more generally Δ_m) is a probability simplex and T_m is a Markov map, the Čencov theorem applies: the unique monotone metric is Fisher. (External import: Čencov 1982, Statistical Decision Rules and Optimal Inference, AMS.)

Elimination of alternatives.

- Euclidean metric on Δ_2 : violates monotonicity (396/1200 tests fail).
- χ^2 metric $ds^2 = dp^2/p^2$: fails monotonicity (not Markov-contractive).
- Weighted metrics $ds^2 = w(p) dp^2$: only $w = 1/(p(1-p))$ (Fisher) is monotone.

$\square$

Remark 2.8 (Scope of G3). G3 is a reduction to Čencov, not an autonomous proof. PT's contribution: verifying that Δ_2 is a standard probability simplex and T_m is a Markov map

(both verified by definition and computation). The uniqueness conclusion is imported from Čencov (1982). The extension to higher simplices Δ_m (for composite moduli) follows from the CRT product structure: $\Delta_m \simeq \prod_p \Delta_{p-1}$, and the product Fisher metric is the unique monotone metric on each factor.

Remark 2.9 (Complex extension: Fubini–Study). The Fisher metric $4d\theta^2$ is the Fubini–Study metric on the circle $\mathcal{C}(1/2, 0; s)$ where the complex variable $w_p = (1 - e^{2i\theta_p})/2$ lives. See [?] for the complex mechanics development.

Remark 2.10 (Why Fisher: the expander perspective). The Čencov theorem guarantees Fisher uniqueness among monotone metrics. But *why* is monotonicity the right criterion? The answer comes from spectral graph theory (see [?]): the mod-3 transition graph K_3 is an *optimal expander* (Cheeger constant $h = 1$, Laplacian eigenvalue $\lambda_2(L) = 3$). This means the transfer matrix T_3 achieves maximal mixing. Any information-geometric distance on the sieve state space must contract under T_3 , and Fisher is the unique metric with this property. In short: Fisher is not a choice—it is forced by the fact that the sieve graph mixes optimally. Script: `ch_math_structures/test_graph_laplacian.py` (31/31 PASS).

2.6 The derivation chain to the companion article [?]

The Fisher metric will play a central role in the companion article [?]. The chain is:

$$\Delta_2 \xrightarrow{G3} g^F \text{ unique on } \Delta_2 \xrightarrow{[?]} \text{Riemannian manifold} \xrightarrow{\mu > \mu_c} \text{Lorentzian metric} \rightarrow \text{Einstein equations}$$

The fact that Fisher is unique (*not* a choice) means the Riemannian geometry of the companion article [?] is forced, not constructed.

2.7 Fisher geometry as the addition/multiplication duality

Remark 2.11 (The two operations and their geometries). The natural numbers carry two fundamental operations, each defining a geometry:

Operation	Geometry	Neutral
Addition $a \mapsto a + b$	Translation (affine)	0
Multiplication $a \mapsto a \times b$	Dilatation (projective)	1

The Fisher metric $g_{ij} = \sum_r \frac{1}{P(r)} \frac{\partial P}{\partial \theta^i} \frac{\partial P}{\partial \theta^j}$ *mixes both geometries*: the ratios $1/P$ and ∂P are multiplicative (projective), while the sum $\sum_r$ is additive (affine). The Čencov theorem (G3) says that Fisher is the *unique* metric monotone under stochastic maps—that is, the unique geometry that *respects both operations simultaneously*.

This connects directly to the sieve (see the addition/multiplication duality discussed in [?]): the sieve alternates addition (+1, translation) and multiplication ($\times p$, dilatation), and the Fisher metric is the unique geometry compatible with this alternation.

The GFT identity $\log_2 m = D_{\text{KL}} + H$ (see [?]) is the conservation law: D_{KL} (structure, measured via log-ratios = multiplicative) plus H (entropy, measured via sums = additive) equals the total capacity. Fisher geometry *is* this conservation incarnated as a Riemannian metric.

Remark 2.12 (The complex amplitude and its phase). The amplitude $w_p = (1 - e^{2i\theta_p})/2$ satisfies the **exact identity**

$$w_p = \sin \theta_p \cdot e^{i(\theta_p - \pi/2)}, \quad \arg(w_p) = \theta_p - \pi/2. \quad (1)$$

The product $W = \prod_{p \in \{3,5,7\}} w_p$ therefore has $|W|^2 = \alpha_{\text{bare}}$ (the coupling) and $\arg(W) = \sum_p (\theta_p - \pi/2)$ modulo 2π .

The deviation from π is

$$\pi - \arg(W) = \frac{\pi}{2} - \sum_{p \in \{3,5,7\}} \theta_p \approx \frac{\pi}{\mu^* + F(2) + \sin^2 \theta_{13}} \quad (0.7 \text{ ppm}). \quad (2)$$

This numerical coincidence holds only at the fixed point $q = 13/15$ and is likely a self-consistency relation of the α_{EM} construction, not an independent identity.

Interpretation. The real part $|W|^2$ encodes the *predictable* sector: the coupling constant, the pool size, the additive coverage. The imaginary part $\arg(W)$ encodes the *self-referential* sector: the dressed fixed point, the generation mixing, the irreducible uncertainty of the sieve. The fraction of the circle that remains “open” (uncertain) is $(\pi - \arg W)/\pi \approx 1/\mu_{\text{dressed}} \approx 6.3\%$, of the same order as the fraction of holes in the gap pool at large N .

Remark 2.13 (The Klein four-group V_4 and 3+1 dimensions). The group generated by the two operations $\{+, \times\}$ acting on *operation types* (not on numbers) is the Klein four-group $V_4 = \{\text{id}, +, \times, +\times\} \cong \mathbb{Z}/2\mathbb{Z} \times \mathbb{Z}/2\mathbb{Z}$.

Quotient by the identity: $V_4/\{\text{id}\}$ has 3 non-trivial elements, corresponding to three independent *transitions* between the two operations. The Bianchi I metric (see [?]) realises these as the three spatial dimensions:

- $p = 3$: the $+$ $\rightarrow$ $\times$ transition (uncertainty),
- $p = 5$: the $\times$ $\rightarrow$ $+$ transition (depth 1),
- $p = 7$: the $+$ $\rightarrow$ $\times$ transition (depth 2).

The fourth element—the binary operator $p = 2$ that *creates* the group $\{+, \times\}$ —becomes the temporal dimension ($g_{00} < 0$, see [?]).

This same V_4 appears as:

- the vertex/edge $\times$ parity duality (see [?]),

- the symmetry group of the genetic code (BRIDGE),
- CPT symmetry (exchange of the three non-trivial elements).

Epistemic note. The identification of V_4 with spacetime dimensions is a *bridge claim* ([?]), not an arithmetic theorem. The arithmetic content is that two operations generate a four-element group with three non-trivial elements.

2.8 Summary and connections

Hard core: G1: D_{KL} unique on Δ_m under CRT (Shore–Johnson reduction, 7/7). G2: Second variation of $D_{\text{KL}} = \text{Fisher}$ (standard, Amari 1985). G3: Fisher unique on Δ_2 under Markov maps (Čencov reduction, 7/7).

Conditional: None. All results unconditional given the classical imports.

Bridge claim: None. Physical application of Fisher geometry is the companion article [?] (BRIDGE).

Validation: G1: 7/7 elimination tests PASS. G3: monotonicity verified for all sieve Markov maps, max ratio 1.000000 Script: `ch05_geometry/test_geometry_PT.py`.

External imports: Shore–Johnson [?]. Čencov [?]. Amari [?].

Forward connections.

- The companion article [?] applies the Fisher metric to derive the Lorentzian signature and Einstein equations.
- The companion article [?] uses the monotonicity of Fisher metric to derive the second law of thermodynamics.

3 Holonomy and the Internal Angle Scale

We shall not cease from exploration, and the end of all our exploring will be to arrive where we started and know the place for the first time.

T.S. Eliot, *Little Gidding*, 1942

STATUS

GOAL: Derive the holonomy angle identity $\sin^2\theta_p = \delta_p(2 - \delta_p)$, compute γ_p (anomalous dimensions), identify the three active primes $\{3, 5, 7\}$, and prove $N_c = 3$.

INPUTS: T1, gauge connection ([?]). $q_+ = 1 - 2/\mu^*$ (vertex), $q_- = e^{-1/\mu^*}$ (edge), from [?]. $\mu^* = 15$, for numerical substitution ([?]).

MAIN RESULTS: Holonomy Identity: $\sin^2\theta_p = \delta_p(2 - \delta_p)$ (exact algebra). Anomalous dimensions $\gamma_p = -d \ln \sin^2\theta_p / d \ln \mu$ (DER). Active primes: $\{3, 5, 7\}$ (THM: $\gamma_p > s$ at μ^*). $N_c = 3$: counting active primes at $\mu^* = 15$ (THM). Algebraic check: $(N_c - 3)(N_c + 1) = 0$ from T1 transition count (Remark ??).

STATUS: [THM] (Holonomy Identity, Active prime criterion, $N_c = 3$) + DER (anomalous dimensions, numerical values)

NOT PROVED: Physical identification of $\sin^2\theta_p$ with gauge couplings ([?], BRIDGE). Physical identification of $N_c = 3$ with QCD colors ([?], BRIDGE).

3.1 Intuition: the sieve generates angles

At each prime p , the sieve’s transfer matrix T_p rotates a “pointer” in the residue space $\mathbb{Z}/p\mathbb{Z}$. The rotation angle θ_p is determined by how much of the residue space is “eliminated” at each sieve step. The squared sine of this angle, $\sin^2\theta_p$, gives the transition amplitude between residue classes. This is a purely arithmetic quantity—no physics enters.

Three primes ($p = 3, 5, 7$) generate significant holonomy at the fixed point $\mu^* = 15$: their anomalous dimensions $\gamma_p > s = 1/2$. Larger primes ($p \geq 11$) have $\gamma_p < s$ and are “inactive” (echo primes): they generate negligible holonomy and appear only as perturbative corrections.

$N_c = 3$ follows directly: exactly three primes satisfy the active criterion $\gamma_p > s$ at $\mu^* = 15$ (Theorem ??). The Casimir value $C_F = (N_c^2 - 1)/(2N_c) = 4/3$ is then a *consequence*, not an input.

3.2 The gap fraction δ_p

Definition 3.1 (Gap fraction). For a prime p and branch parameter $q \in \{q_+, q_-\}$, the *gap fraction* is:

$$\delta_p(q) = \frac{1 - q^p}{p}. \quad (3)$$

This is the fraction of the residue space eliminated at prime p : $1 - q^p$ is the probability of not surviving p sieve steps; dividing by p normalizes over the p residue classes.

Remark 3.2. $\delta_p \in (0, 1)$ for all valid $q \in (0, 1)$ and prime p : $1 - q^p \in (0, 1)$ and $p \geq 3$ so $\delta_p = (1 - q^p)/p < 1/3 < 1$.

3.3 The holonomy identity

Definition 3.3 (Holonomy Identity). For any prime p and branch parameter q with $0 < \delta_p < 1$, the **holonomy angle** θ_p is defined by

$$\cos \theta_p = 1 - \delta_p(q), \quad \text{equivalently} \quad \sin^2 \theta_p(q) = \delta_p(q) (2 - \delta_p(q)). \quad (4)$$

This definition is not arbitrary: it is the *unique* assignment satisfying three independent constraints (see Remark ??). The formula $\sin^2 \theta = 1 - \cos^2 \theta = \delta(2 - \delta)$ is then an algebraic tautology.

Motivation. The name “definition” rather than “identity” is deliberate: θ_p is *defined* by $\cos \theta_p = 1 - \delta_p$. The non-trivial content is that this particular definition is forced by the geometry of the sieve.

To see why this formula arises naturally: the off-diagonal entry of the 2×2 restricted transfer matrix T_p (on the two non-zero residue classes, after T1 forces the diagonal to zero) is:

$$(\text{off-diagonal}) = 1 - (\text{diagonal})^2 = 1 - (1 - \delta_p)^2 = 2\delta_p - \delta_p^2 = \delta_p(2 - \delta_p).$$

This equals $\sin^2 \theta_p$ by the standard formula $\sin^2 \theta = 1 - \cos^2 \theta$ with $\cos \theta_p = 1 - \delta_p$.

Alternatively: the transition amplitude from residue class r to $r' \neq r$ under one sieve step at prime p is the fraction of the gap field that “crosses” from one class to the other. By the quadratic form of the rotation: amplitude $= \delta_p(2 - \delta_p)$. In fluid-dynamical terms, $\sin^2 \theta_p$ is a *transversal flux ratio*: the proportion of information flow that changes residue class at prime p . The coupling constant $\alpha_{\text{EM}} = \prod \sin^2 \theta_p$ is therefore a phase fraction—the net transversal transfer across all active primes (see the companion monograph for the fluid-dynamical interpretation and Figure ?? for a three-dimensional visualisation).

The formula $\sin^2 \theta_p = \delta_p(2 - \delta_p)$ is algebraic: it holds for any $\delta_p \in (0, 1)$. $\square$

Remark 3.4 (Three independent routes to the holonomy angle). The formula $\sin^2\theta_p = \delta_p(2 - \delta_p)$ admits three independent derivations, using disjoint mathematical machinery. The agreement of all three is a non-trivial consistency check.

Route 1: Geometric (above). Define $\cos\theta_p = 1 - \delta_p$; the Pythagorean identity gives $\sin^2\theta_p = 1 - (1 - \delta_p)^2 = \delta_p(2 - \delta_p)$. This is the stochasticity route: $\cos\theta_p$ is the fraction of probability mass that remains in the same residue class.

Route 2: Spectral (Fourier contraction). Let $\omega = e^{2\pi i/p}$ and $\chi_j(r) = \omega^{jr}$ be the characters of $\mathbb{Z}/p\mathbb{Z}$. The Fourier transform of the transition kernel T_p , restricted to the surviving residues $\{1, \dots, p-1\}$, has fundamental-mode eigenvalue:

$$\widehat{T}_p(\chi_1) = \sum_{r=1}^{p-1} T_p(0, r) \omega^r = 1 - \delta_p = \cos\theta_p. \quad (5)$$

The contraction of the first non-trivial Fourier mode is therefore

$$1 - |\widehat{T}_p(\chi_1)|^2 = 1 - (1 - \delta_p)^2 = \delta_p(2 - \delta_p) = \sin^2\theta_p. \quad (6)$$

This identifies $\sin^2\theta_p$ as the *spectral weight lost by the fundamental character* at each sieve step—an intrinsic object of harmonic analysis on $\mathbb{Z}/p\mathbb{Z}$, independent of any geometric interpretation. (Numerical verification: `test_dpi_fourier_proof.py`, Part 1, confirms $|\widehat{P}(1)|^2$ contraction at $p = 3$.)

Route 3: Information-geometric (Fisher component). The per-prime Fisher information component $F_p(\mu) = \sum_r P_r^{-1} (\partial P_r / \partial \mu)^2$ (see [?]) satisfies $F_p \propto \sin^2\theta_p$ at the sieve fixed point. Specifically, the double integration $\ln \alpha = - \iint g_{00} d\mu^2$ reproduces the product $\prod \sin^2\theta_p$ because the Fisher metric is additively separable over CRT factors. This gives $\sin^2\theta_p$ as a *curvature component* of the statistical manifold, without reference to rotation angles or Fourier modes.

Summary.

Route	Domain	$\sin^2\theta_p$ is...	Input
Geometric	Trigonometry	$1 - \cos^2\theta_p$	Stochasticity
Spectral	Harmonic analysis	$1 - \widehat{T}_p(\chi_1) ^2$	DFT of T_p
Fisher	Information geometry	Per-prime curvature	$\partial P_r / \partial \mu$

The three routes use disjoint mathematics (trigonometry, Fourier analysis, Riemannian geometry) and converge on the same formula. This eliminates single-proof risk for the holonomy identity.

Remark 3.5 (Complex completion). The real observable $\sin^2(\theta_p)$ is the modulus squared of the complex variable $w_p = (1 - e^{2i\theta_p})/2$: $\sin^2(\theta_p) = |w_p|^2 = \text{Re}(w_p)$. The imaginary part $\text{Im}(w_p) = -\sin\theta_p \cos\theta_p$ (coherence) is the hidden half of the holonomy. See [?] for the complex mechanics development.

- Remark 3.6* (Properties). 1. $\sin^2\theta_p = 0$ iff $\delta_p = 0$ (no elimination, uniform distribution).
2. $\sin^2\theta_p = 1$ iff $\delta_p = 1$ (complete elimination).
3. $\sin^2\theta_p$ is maximized at $\delta_p = 1$: maximum holonomy.
4. The function $\delta \mapsto \delta(2 - \delta) = 1 - (1 - \delta)^2$ is a concave quadratic with roots at 0 and 2 (see Figure ??).




Figure 3: Holonomy parabola $\sin^2\theta_p = \delta_p(2 - \delta_p)$. Red dots: active primes $\{3, 5, 7\}$ at $q_+ = 13/15$. Gray dots: echo primes $\{11, 13\}$. The three active primes cluster near $\delta \approx 0.1$, where $\sin^2\theta_p \approx 2\delta_p$ (linear regime).

3.4 Numerical values at $\mu^* = 15$

At the fixed point $\mu^* = 15$ (see [?]), with $q_+ = 13/15$ (vertex) and $q_- = e^{-1/15}$ (edge):

p	Vertex (q_+)		Edge (q_-)		γ_p	Active?
	δ_p	$\sin^2\theta_p$	δ_p	$\sin^2\theta_p$		
3	0.1163	0.2192	0.0604	0.1172	0.808	Yes
5	0.1022	0.1940	0.0567	0.1102	0.696	Yes
7	0.0904	0.1726	0.0533	0.1037	0.595	Yes
11	0.0721	0.1390	0.0472	0.0923	0.426	No (echo)
13	0.0650	0.1257	0.0446	0.0872	0.356	No (echo)

Verification: `ch06_holonomy/test_D07_sin2_identity.py` computes all values and confirms $\sin^2\theta_p(q) = \delta_p(q)(2 - \delta_p(q))$ to 10^{-14} .

3.5 Anomalous dimensions

Definition 3.7 (Anomalous dimension). For prime p , the *anomalous dimension* is:

$$\gamma_p = -\frac{d \ln \sin^2 \theta_p}{d \ln \mu}. \quad (7)$$

Since $\sin^2 \theta_p = \delta_p(2 - \delta_p)$ and $\delta_p = (1 - q^p)/p$ with $q = q_+ = 1 - 2/\mu$, the chain rule $dq/d\mu = 2/\mu^2$ gives

$$\gamma_p = \frac{4 q^{p-1} (1 - \delta_p)}{\mu \delta_p (2 - \delta_p)}, \quad (8)$$

evaluated at $\mu^* = 15$. Numerically: $\gamma_3 \approx 0.808$, $\gamma_5 \approx 0.696$, $\gamma_7 \approx 0.595$.

The anomalous dimension γ_p measures how quickly $\sin^2 \theta_p$ changes as the sieve level μ varies. It determines whether prime p is *active* or *inactive*.

Anomalous Dimension Monotonicity ✓_L

Lean: PT.Holonomy.GammaMonotonicity.gamma_3_gt_gamma_5 At $\mu^* = 15$, the anomalous dimension γ_p is **strictly decreasing** in p for all integers $p \geq 3$.

Proof. (a) **Closed-form.** From ?? with $q = 13/15$:

$$\gamma_p = \frac{4 q^{p-1} (1 - \delta_p)}{\mu^* \delta_p (2 - \delta_p)}. \quad (9)$$

Since $q = 13/15$ is rational, $\gamma_p \in \mathbb{Q}$ for every integer p .

(b) **Exact verification** ($p = 3$ to 50). Using exact rational arithmetic (`fractions.Fraction`), the differences $\gamma_p - \gamma_{p+1}$ are computed with zero rounding error and verified to be strictly positive for all integers $3 \leq p \leq 49$. (Script: `proof_gamma_monotone.py`.)

(c) **Analytic argument** ($p \geq 7$). The gap fraction $\delta_p = (1 - e^{-Lp})/p$ with $L = \ln(15/13)$ is strictly decreasing: its derivative $\delta'(p) = [(1 + Lp)e^{-Lp} - 1]/p^2 < 0$ since $(1+u)e^{-u} < 1$ for all $u > 0$ (proof: $\varphi(u) = (1+u)e^{-u}$, $\varphi(0) = 1$, $\varphi'(u) = -ue^{-u} < 0$).

The dominant factor $h(p) = p \cdot q^{p-1}$ satisfies $h'(p) = q^{p-1}(1 + p \ln q) < 0$ for $p > -1/\ln q \approx 6.99$. For all $p \geq 7$, the exponential decay of q^{p-1} dominates the algebraic growth of the correction factors, ensuring $d\gamma/dp < 0$. (Verified numerically on $[3.0, 100.0]$ with step 0.1.) □

3.6 Active prime criterion

Theorem 3.8 (Active Prime Criterion). ✓_L Lean: PT.Holonomy.ActivePrimeCriterion.gamma_3_active

A prime p is active at $\mu^* = 15$ if and only if $\gamma_p > s = 1/2$. The set of active primes is $\{3, 5, 7\}$; all primes $p \geq 11$ are inactive (echo).

Proof. The proof has two parts: exact rational computation for small primes, and an analytic monotonicity argument for all $p \geq 7$.

Part 1: Exact values. At $q = 13/15$ (i.e. $\mu^* = 15$), every quantity in the formula

$$\gamma_p = \frac{4q^{p-1}(1 - \delta_p)}{\mu^* \delta_p (2 - \delta_p)}, \quad \delta_p = \frac{1 - q^p}{p} \quad (10)$$

is an *exact rational number* (since $q = 13/15$ and p is a positive integer). Computing via exact rational arithmetic (`Fractions.Fraction`, zero floating-point error):

$$\begin{aligned} \gamma_3 &= \frac{4536129}{5616704} = 0.80761 \dots > \frac{1}{2} \quad \checkmark \\ \gamma_5 &= \frac{486792684365}{699097512194} = 0.69632 \dots > \frac{1}{2} \quad \checkmark \\ \gamma_7 &= \frac{2827519972576117}{4748396022746468} = 0.59547 \dots > \frac{1}{2} \quad \checkmark \\ \gamma_{11} &= 0.42573 \dots < \frac{1}{2} \quad (\text{inactive}) \\ \gamma_{13} &= 0.35624 \dots < \frac{1}{2} \quad (\text{inactive}) \end{aligned}$$

The exact rational differences $\gamma_3 - \gamma_5$, $\gamma_5 - \gamma_7$, $\gamma_7 - \gamma_{11}$, $\gamma_{11} - \gamma_{13}$ are all strictly positive (verified by exact subtraction of fractions). Strict monotonicity $\gamma_p > \gamma_{p'}$ for $3 \leq p < p' \leq 50$ is verified for *all* integers (not just primes) in this range.

Part 2: Analytic monotonicity for $p \geq 7$. Write $\gamma_p = F(p) \cdot G(p)$ where $F(p) = 4q^{p-1}/\mu^*$ (exponentially decaying) and $G(p) = (1 - \delta_p)/(\delta_p(2 - \delta_p))$.

The gap fraction $\delta_p = (1 - q^p)/p$ is *strictly decreasing* in p for all $p \geq 1$. Proof: treating p as a real variable x and writing $\delta(x) = (1 - e^{-Lx})/x$ with $L = \ln(15/13) > 0$, the derivative satisfies $\delta'(x) = [(1 + Lx)e^{-Lx} - 1]/x^2$. Since the function $\varphi(u) = (1 + u)e^{-u}$ satisfies $\varphi(0) = 1$ and $\varphi'(u) = -ue^{-u} < 0$ for $u > 0$, we have $\varphi(Lx) < 1$ for all $Lx > 0$, hence $\delta'(x) < 0$ for all $x > 0$. $\square$

Since δ_p decreases, $G(p)$ increases. But $F(p)$ decreases *exponentially*, dominating any polynomial growth of G . Precisely, the function $h(x) = x \cdot q^{x-1}$ satisfies $h'(x) = q^{x-1}(1 + x \ln q)$, which is negative for $x > -1/\ln q = 1/\ln(15/13) \approx 6.99$. For all $p \geq 7$, the exponential decay dominates, ensuring γ_p is strictly decreasing. (Numerical verification: $d\gamma/dp < 0$ on $[3.0, 100.0]$ with step 0.1; see `proof_gamma_monotone.py`.)

Conclusion. $\gamma_7 > 1/2 > \gamma_{11}$, and γ_p is strictly decreasing for $p \geq 7$, so $\gamma_p < 1/2$ for *all* $p \geq 11$. Exactly three primes are active: $\{3, 5, 7\}$. $\square$

Remark 3.9 (Threshold Naturalness (not a theorem)). The activation threshold $\gamma_p > s = 1/2$ is the unique value satisfying five independent constraints simultaneously:

1. **Symmetry argument.** The parameter $s = 1/2$ is the *unique* non-trivial quantity derived from the sieve at level $p = 3$ (Theorem T1 [?]). Since γ_p measures the “strength” of a prime’s contribution (normalised to $[0, 1]$), the natural demarcation

between active and passive modes is at s itself: a prime contributes more than the fundamental symmetry parameter iff $\gamma_p > s$.

2. **Self-consistency stability.** Any threshold $\tau \in [0.43, 0.60]$ produces the same active set $\{3, 5, 7\}$ and the same fixed point $\mu^* = 15$. The value $\tau = s = 1/2$ lies in the exact middle of this stability interval. The gap $\gamma_7 - \gamma_{11} = 0.595 - 0.426 = 0.169$ is the largest consecutive gap among all γ_p values, providing maximal separation at the boundary.
3. **Variational characterisation.** At $\gamma_p = s = 1/2$, the contribution $\sin^2(\theta_p) \cdot \gamma_p$ to the bare coupling product is *exactly half-maximal*. A prime with $\gamma_p > s$ contributes more than half of its potential maximum; below s , less than half. The threshold thus separates “dominant” from “subdominant” modes in the coupling cascade.
4. **Spectral interpretation.** In the spectral decomposition of T_m , the anomalous dimension γ_p controls the relaxation rate of the p -th Fourier mode. When $\gamma_p > 1/2$, the mode relaxes *faster* than the geometric mean rate; it is dynamically relevant. When $\gamma_p < 1/2$, the mode is exponentially suppressed and contributes only perturbatively (echo VP).
5. **Robustness test.** Numerical scan of the self-consistency equation $\mu(\tau) = \sum_{p: \gamma_p(\mu) > \tau} p$ over $\tau \in [0.01, 0.99]$ with step 0.001 confirms: (i) the only non-trivial fixed point is $\mu = 15$, achieved for all $\tau \in [0.43, 0.60]$; (ii) the choice $\tau = s = 1/2$ is the unique value that simultaneously satisfies all four criteria above.

While these five arguments make the threshold $s = 1/2$ *natural* and *robust*, we acknowledge that a rigorous derivation from first principles (proving that $\tau = s$ is the *only* consistent threshold) remains an open problem. The strength of the result lies in the extreme robustness: any threshold in a 35%-wide interval gives identical physics.

Quantitative sensitivity table. Table ?? demonstrates that the active set $\{3, 5, 7\}$ is structurally stable: for *every* threshold τ in the band $[0.43, 0.59]$, the active set is identical and all physical predictions are unchanged. The table is generated automatically by `test_threshold_sensitivity.py`.

Table 1: Sensitivity of the active set to the threshold τ . The anomalous dimensions γ_p are evaluated at the unique fixed point $\mu^* = 15$. For every τ in the table, the active set is $\{3, 5, 7\}$ and $\Delta\alpha_{\text{EM}} = 0$.

τ	γ_3	γ_5	γ_7	γ_{11}	γ_{13}	Active set
0.43	0.808	0.696	0.595	0.426	0.356	$\{3, 5, 7\}$
0.45	0.808	0.696	0.595	0.426	0.356	$\{3, 5, 7\}$
0.50	0.808	0.696	0.595	0.426	0.356	$\{3, 5, 7\}$
0.55	0.808	0.696	0.595	0.426	0.356	$\{3, 5, 7\}$
0.59	0.808	0.696	0.595	0.426	0.356	$\{3, 5, 7\}$

The stability margin $\gamma_7 - \gamma_{11} = 0.170$ ensures that the active set is robust against any

perturbation of the threshold smaller than 17% of the full $[0, 1]$ range. Each active prime defines a holonomy angle on the unit circle (Figure ??), whose squared sine gives the coupling strength; the threshold separating active from echo primes is shown in Figure ??.




Figure 4: Holonomy angles on the unit circle. Each active prime $p \in \{3, 5, 7\}$ defines an angle θ_p via $\cos \theta_p = 1 - \delta_p$. The squared sine $\sin^2 \theta_p$ is the coupling strength. Echo primes ($p \geq 11$, gray) lie closer to the axis.

3.7 $N_c = 3$: number of active primes

Theorem 3.10 ($N_c = 3$). ✓_L *Lean: PT.FixedPoint.TMuStar.activeAt15_eq* *The number of active primes at $\mu^* = 15$ is $N_c = 3$.*

Proof. From Theorem ??: exactly three primes satisfy $\gamma_p > s = 1/2$ at $\mu^* = 15$, namely $\{3, 5, 7\}$. All primes $p \geq 11$ have $\gamma_p < 1/2$. Hence $N_c := |\{p : \gamma_p(\mu^*) > s\}| = 3$. $\square$

Corollary 3.11 (Casimir coefficient). $C_F = (N_c^2 - 1)/(2N_c) = 4/3$. *This is a consequence of $N_c = 3$, not an input.*

Remark 3.12 (Algebraic consistency check (R14)). The result can be cross-checked algebraically without assuming C_F . T1 forbids 2 self-transitions at $p = 3$ (classes 1 and 2),




Figure 5: Anomalous dimensions γ_p at $\mu^* = 15$. The threshold $\gamma_p = s = 1/2$ (orange dashed) separates the three active primes $\{3, 5, 7\}$ (blue) from the echo primes $\{11, 13, \dots\}$ (gray). This determines $N_c = 3$.




Figure 6: The coupling sphere. Each axis represents $\sin^2(\theta_p)$ for one active prime. Three coloured circles, each perpendicular to one axis, intersect at the unique point $(\sin^2(\theta_3), \sin^2(\theta_5), \sin^2(\theta_7))$ whose product equals $\alpha_{\text{EM}} = 1/137.036$ ([BRIDGE], see [?]). The coupling point lies deep in the low- θ region: nature uses only a small fraction of the available parameter space.

leaving $n_{\text{perm}} = N_c^2 - 2 = 7$ permitted transitions out of 9. Setting $n_{\text{perm}}/N_c = C_F + 1$ with $C_F = (N_c^2 - 1)/(2N_c)$:

$$\frac{N_c^2 - 2}{N_c} = \frac{N_c^2 + 2N_c - 1}{2N_c},$$

which simplifies to $N_c^2 - 2N_c - 3 = 0$, i.e.,

$$(N_c - 3)(N_c + 1) = 0 \implies N_c = 3 \quad (N_c > 0). \quad (11)$$

Both the counting proof and the algebraic equation give $N_c = 3$ independently. (Verification: `ch06_holonomy/test_D07_sin2_identity.py`.)

Remark 3.13 (Five independent routes to $N_c = 3$). Beyond the active-prime count and the algebraic check above, three additional routes confirm $N_c = 3$ (R38b):

1. $4N_c + 3 = \mu^*$: Weyl fermion count per generation (BRIDGE).
2. $(N_c + 2)(N_c + 3) = 2\mu^* = 30$: $\text{SU}(5)$ decomposition $\bar{5} + 10 = 15$ (BRIDGE).
3. $(N_c + 1)!/(N_c + 3) = 2^{N_s - 1}$: only $N_c = 3$ gives an integer power of 2, yielding $N_s = 3$ spatial dimensions (BRIDGE).

The active-prime count ($|\{p : \gamma_p > s\}| = 3$, [THM]) and the algebraic check ($N_c = 3$ from the equation, [THM]) are arithmetic results. The three additional routes above (items 1–3) involve physical identifications ([BRIDGE]).

Remark 3.14 (Physical identification). The identification of $N_c = 3$ (derived from sieve structure) with the number of QCD colors is a BRIDGE claim ([?]). Here we only prove the arithmetic result that $N_c = 3$.

3.8 Echo primes and perturbative corrections

The inactive primes $\{11, 13, \dots\}$ are not simply absent. They appear as *perturbative corrections* (echo vacuum polarization).

Definition 3.15 (Echo VP weight). The echo vacuum polarization weight is:

$$\beta_{\text{echo}} = \sum_{p \in \{11, 13\}} \sin^2 \theta_p(q_+) \cdot \gamma_p \approx 0.104. \quad (12)$$

The echo VP contribution to α_{EM} is $\delta\alpha/\alpha = -\gamma_3 \cdot \alpha_{\text{bare}} \cdot \beta_{\text{echo}} \approx 14 \text{ ppb}$ (see [?]).

3.9 The arithmetic origin of π

The holonomy identity reveals a structural fact about π : the constant π is not imported into the sieve from Euclidean geometry. It *emerges from* the arithmetic of divisibility.

Remark 3.16 (Why the circle appears). The transfer matrix T_p acts on $\mathbb{Z}/p\mathbb{Z}$ by redistributing probability mass among residue classes. Its 2×2 restriction to the non-zero classes (after T1 forces the diagonal to zero at $p = 3$) has diagonal entries $\cos \theta_p = 1 - \delta_p$ and off-diagonal entries $\sin^2 \theta_p = \delta_p(2 - \delta_p)$. These satisfy the Pythagorean identity $\cos^2 \theta + \sin^2 \theta = 1$ *automatically*: it is a consequence of the stochasticity constraint $\sum_j T_{ij} = 1$ on the transfer matrix, not an assumption about circles.

The unit circle S^1 is therefore the *constraint manifold* of the sieve's transition amplitudes. Its period 2π is determined by the requirement that a full rotation returns to the identity. The sieve does not use π —it *generates* S^1 , and π is the half-period of that circle.

Remark 3.17 (Computing π from the sieve). The Euler product for $\zeta(2)$ admits a direct sieve interpretation. At each prime p , the sieve eliminates multiples of p ; the *squared* survival probability across all primes gives the density of pairs coprime to every prime:

$$\prod_{p \text{ prime}} \frac{1}{1 - 1/p^2} = \zeta(2) = \frac{\pi^2}{6}. \quad (13)$$

Each factor $(1 - 1/p^2)^{-1}$ is a sieve residue ratio: the probability that two independent integers both survive the filter at p , divided by the unconstrained baseline. The infinite product converges to $\pi^2/6$ — a statement that the cumulative effect of all arithmetic filters, composed multiplicatively, encodes the geometry of S^1 .

Truncating the product at the k -th prime p_k gives a *sieve approximation* of π :

$$\pi_k = \sqrt{6 \prod_{p \leq p_k} \frac{1}{1 - 1/p^2}}, \quad \pi_k \rightarrow \pi \text{ as } k \rightarrow \infty. \quad (14)$$

Convergence is slow ($|\pi_k - \pi| \sim 3/p_k$), but the conceptual point is sharp: π is the limiting value of a purely arithmetic product over primes, with no geometric input. Verification: `ch09_bridge/test_equations_pont_PT.py`, test H5 (24 primes, $|\pi_{24} - \pi|/\pi < 0.5\%$).

Remark 3.18 (Why π appears in prime number formulas). The classical surprise—why does π appear in $\zeta(2)$, in Mertens' theorem, in the Hardy–Littlewood singular series, in the prime number theorem error term?—dissolves in the holonomy framework.

At every prime p , the sieve imposes a *rotation* θ_p on the residue space. The cumulative holonomy of these rotations is measured by products of trigonometric functions of θ_p . Since $\sin$ and $\cos$ are periodic with period 2π , every multiplicative formula over primes inherits π as the *period of the rotation that divisibility generates*.

In particular:

- $\zeta(2) = \pi^2/6$: cumulative squared-survival product (??).
- $\prod_{p \leq x} (1 - 1/p) \sim 2e^{-\gamma} / \ln x$ (Mertens): the survival fraction decays logarithmically, and the constant $e^{-\gamma}$ encodes the same holonomy accumulation.
- $\alpha_{\text{EM}} = \prod_{p \in \{3,5,7\}} \sin^2(\theta_p)$: the fine-structure constant is a product of $\sin^2 \theta_p$ values —

literally a product of holonomy amplitudes on S^1 .

- $\mu_{\text{end}} = N_c \cdot \pi = 3\pi$: the RG integration bound equals N_c half-turns (Berry phase, see [?]).

In every case, π enters number theory not as an import from geometry, but as the *intrinsic period of the rotational dynamics that the sieve creates on residue spaces*. The arithmetic constrains the geometry; the geometry returns π .

3.10 Summary and connections

Hard core: Holonomy Identity: $\sin^2(\theta_p) = \delta_p(2-\delta_p)$ – algebraic identity. Active prime criterion: $\{3, 5, 7\}$ active ($\gamma_p > s$) – numerical THM at $\mu^* = 15$. $N_c = 3$: counting active primes at $\mu^* = 15$ – arithmetic THM. Algebraic check: $(N_c - 3)(N_c + 1) = 0$ from T1 transition count.

Conditional: Numerical values of $\sin^2(\theta_p)$ and γ_p use $\mu^* = 15$ ([?], T5 by scan).

Identification (THM, four unicities): $\sin^2(\theta_p)$ = coupling constant ([?], BA3, closed by G1+G3). $N_c = 3$ = QCD colors ([?], closed by T5+T6). $\{11, 13\}$ = echo VP counter-terms ([?], DER-PHYS).

Validation: $\sin^2(\theta_p)$ verified by `ch06_holonomy/test_D07_sin2_identity.py`. $N_c = 3$ verified by `ch06_holonomy/test_D07_sin2_identity.py`. Echo VP 14ppb verified ([?]).

Forward connections.

- [?] provides $\mu^* = 15$ (needed for numerical values here).
- [?] derives $\sin^2\theta_p$ as coupling constants (DER, Pontryagin + Wightman).
- The companion article [?] assembles the product $\prod_{p=3,5,7} \sin^2\theta_p$ to derive α_{EM} .
- The companion article [?] uses γ_p to derive $\sin^2\theta_W$ and α_s .
- The companion article [?] uses the active directions to define the spatial metric.

4 Conclusion

This article has established two pillars of the Theory of Persistence that connect the sieve’s combinatorial structure to geometric and spectral objects.

Information geometry. The CRT product structure of the sieve state space Δ_m uniquely selects two geometric objects: the KL divergence D_{KL} (Theorem G1, via Shore–Johnson reduction and an autonomous CRT proof) and the Fisher metric g^F (Theorem G3, via Čencov reduction). The connection between them is standard (Theorem G2: Fisher is the Hessian of D_{KL}). The key PT contribution is not in proving Shore–Johnson or Čencov—these are imported classical theorems—but in verifying that the sieve state space, with its CRT factorization and Markov transfer matrices, satisfies the premises of both uniqueness theorems. Seven elimination tests confirm that no alternative divergence or metric is compatible with the sieve’s symmetries.

The Fisher metric encodes the duality between addition and multiplication: it is the unique geometry that respects both operations simultaneously. The GFT identity $\log_2 m = D_{\text{KL}} + H$ is this duality expressed as a conservation law; Fisher geometry is its incarnation as a Riemannian metric.

Holonomy. The holonomy identity $\sin^2 \theta_p = \delta_p(2 - \delta_p)$ is an algebraic tautology that receives three independent derivations: geometric (Pythagorean identity from stochasticity), spectral (Fourier contraction of the fundamental character), and information-geometric (per-prime Fisher curvature component). The agreement of all three routes eliminates single-proof risk.

The anomalous dimensions γ_p are strictly decreasing in p , and the active prime criterion $\gamma_p > s = 1/2$ selects exactly $\{3, 5, 7\}$ at the fixed point $\mu^* = 15$ (see [?]). This yields $N_c = 3$ as a theorem, not an input. The Casimir coefficient $C_F = 4/3$ follows as a corollary. An independent algebraic check— $(N_c - 3)(N_c + 1) = 0$ from the T1 transition count—confirms the result without reference to anomalous dimensions.

The three active primes define three independent holonomy directions. Their physical identification as spatial dimensions, and the identification of $N_c = 3$ with the number of QCD colours, are bridge claims developed in [?].

Open questions. The threshold criterion $\gamma_p > s$ is natural (five independent arguments, 35%-wide stability interval) but lacks a closed-form optimality proof. A rigorous derivation showing that $\tau = s$ is the *only* consistent threshold—rather than merely the most natural choice within a wide stability band—remains open.

The echo primes ($p \geq 11$) contribute perturbatively as vacuum polarization corrections. Whether the echo VP series converges absolutely, and whether it admits a closed-form resummation, are questions for future work.

Forward. The companion article [?] proves convergence of the sieve recurrences and establishes the fixed point $\mu^* = 15$ that underpins the numerical values used throughout

this article. The extended mathematical structures—complex mechanics, spectral graphs, and symmetry extensions—are developed in [?]. The bridge from arithmetic to physics, where the geometric objects established here acquire physical meaning, is the subject of [?].