The Theory of Persistence IV: Extended Mathematical Structures

Yan Senez

yan.senez@protonmail.com

Abstract

This is the fourth of five articles presenting the mathematical foundations of the Theory of Persistence (PT). Building on the sieve dynamics and conservation laws [2], the information geometry and holonomy [3], and the convergence theory and fixed point [4], we develop the extended mathematical structures that the Eratosthenes sieve carries beyond the core arithmetic theorems.

Sieve algebra and transforms. The sieve product $*_T$ defines a new algebra of arithmetic functions: non-associative, non-commutative, and without identity—distinct from Dirichlet convolution, Hecke algebras, and $\ell^1(\mathbb{N})$. The persistence transform P_K projects onto the eigenbasis of the transfer matrix T_3; the PT metric d_{PT} is quasi-orthogonal to both the Archimedean and p-adic metrics; the category Cat(Sieve) admits four coherent functors converging to the asymptotic sieve structure. Three new transforms—tensor $\mathcal{T}$ (inter-modular correlations in a $3\times5\times7$ tensor), holonomic $\mathcal{H}$ (generalised Euler product with Fisher metric), and decoherence $\mathcal{D}$ (entropy monotonicity: the second law of the sieve, Theorem H)—extend the classical persistence transform. The gap residue code achieves distance $d = n - 1$ (asymptotically MDS); prediction equals error correction.

Complex mechanics. The real observable $\sin^2\theta_p$ lifts to a complex variable w_p on a circle of radius $s = 1/2$, with $|w_p|^2 = \mathrm{Re}(w_p)$. The holomorphic force $F(w) = i/w - 2i$, Liouville integrability on $\mathbb{T}^n$, an exact uncertainty relation $|p_w|\cdot|w| = 1$, and the QED radiative corrections $3\alpha/(4\pi)$ and $16\pi^2$ emerge from the circle geometry.

Predictive Mathematics. The PM framework quantifies constraint cascades: each prime adds exactly one degree of freedom (Theorem L1, unconditional); the layer dimension follows $\dim(d) = P(d+1) - (2d+1)$ (13 shells verified); the activation spectrum is $\Sigma_k = \binom{4}{k}$ for $k \geq \mathrm{AT}$; and the activation threshold formula in three regimes (echo, construction, asymptotic) proves $\mathrm{AT} = 1$ for all $p \geq 41$.

Convergence extensions. The arithmetic–geometric race provides five alternative routes to T4. The mod-3 transition graph K_3 is an optimal Ramanujan expander whose Ihara zeta function has all non-trivial zeros on the critical circle.

The dihedral group $D_4 = \langle T_3, J \rangle$ organises the spectral structure via the intertwiner J.

Contents

1 Introduction

This article is the fourth instalment (M4) in a five-part series developing the mathematical foundations of the Theory of Persistence (PT). The series is organised as follows:

M1 [2]: The sieve as a dynamical field, uniqueness of Eratosthenes, conservation laws, and the Gap Fundamental Theorem.

M2 [3]: Information geometry, holonomy, and the internal angle scale.

M3 [4]: Convergence, exact recurrences, and the fixed point $\mu^* = 15$.

M4 (this article): Extended mathematical structures: sieve algebra, complex mechanics, Predictive Mathematics, and convergence extensions.

M5 [5]: The bridge from arithmetic to physics.

These five articles together provide a self-contained mathematical foundation of the Theory of Persistence, intended to be read in sequence M1 $\to$ M5.

The companion articles [2–4] established the core arithmetic theory: the sieve dynamics and conservation laws (M1), the canonical information geometry and holonomy angles $\sin^2\theta_p = \delta_p(2-\delta_p)$ (M2), and the convergence to the unique fixed point $\mu^* = 15$ with active primes $\{3, 5, 7\}$ (M3). This article develops the extended mathematical infrastructure that the sieve carries beyond these core theorems.

The material is organised in four sections, corresponding to four distinct mathematical domains.

Section 2: Mathematical structures. The sieve product $*_T$ weights Dirichlet convolution by the transfer matrix T_3, producing a new algebraic structure. The persistence transform P_K, the PT metric d_{PT}, the sieve category, and three new transforms (tensor, holonomic, decoherence) reveal the depth of the sieve's mathematical content. The gap residue code achieves Hamming distance $d = n - 1$ (Theorem, asymptotically MDS): 32 companion scripts validate the full catalogue.

Section 3: Complex mechanics. The lift from the real observable $\sin^2\theta_p$ to the complex variable $w_p = (1 - e^{2i\theta_p})/2$ uncovers hidden structure: all w_p lie on a circle of radius $s = 1/2$ (Circle theorem), the force $F(w) = i/w - 2i$ is holomorphic, angular momentum equals $\sin^2\theta_p$ (Noether charge), and the system is Liouville-integrable on $\mathbb{T}^n$. The QED radiative corrections decompose entirely into PT-native quantities.

Section 4: Predictive Mathematics. The PM framework predicts what becomes structurally inevitable when constraints are stacked on a raw field. Seven theorems from a single algebraic foundation (indicator algebra + CRT): L1 (unit increment), the dimension formula, the activation spectrum, the three-regime AT formula, the partition identity C1, the phantom saturation C2, and the PM–PT corollary (AT = 1 for all $p \geq 41$). Three instantiations (sieve, error-correcting codes, Lucky numbers) sharpen the universal/specific classification.

Section 5: Convergence extensions. The arithmetic–geometric race, spectral graph theory of the sieve (optimal Ramanujan expander, Ihara zeta function), and the D_4 symmetry group deepen and validate the T4 convergence theorem of [4].

All results follow from $s = 1/2$ and the sieve structure established in [2–4]. No physics is invoked; no parameter is fitted.

2 Mathematical Structures of the Sieve

> The purpose of computing is insight, not numbers.
>
> Richard Hamming

STATUS

GOAL: Expose the algebraic, metric, categorical, quantum, and coding-theoretic structures that the sieve carries beyond the physical applications of the companion articles [4, 7].

INPUTS: T_3 (the companion article [2]), GFT (the companion article [2]), Fisher geometry (the companion article [3]), CRT structure (the companion article [2]), spectral bound (the companion article [4]).

MAIN RESULTS:

- New sieve algebra $*_T$: non-associative, non-commutative, no identity (M15).
- Persistence transform P_K: exact spectral decomposition (M16).
- PT metric d_{PT}: quasi-orthogonal to $|\cdot|$ and $|\cdot|_p$ (M17).
- PT numbers $\mathbb{Z}_{\mathrm{PT}}$: profinite enrichment of $\mathbb{Z}$ (M18).
- Category Cat(Sieve): four coherent functors (M19).
- Universal modular extension: structures hold for all primes q (M21).
- Quantum sieve: CPTP channel with arithmetic decoherence (M27).
- CRT tensor network: MPS with bond dimension 3 exact (M28).
- Multi-modular prediction: +61% over mod-3 alone (M30).
- CRT quantum code $[[3, 5.58, 1]]$: prediction = error correction (M31).
- Gap distance code: $d = n - 1$ (THEOREM), asymptotically MDS (M32).
- Tensor transform $\mathcal{T}$: $3{\times}5{\times}7 = 105$ components, rank classifies inter-modular correlations (M33).
- Holonomic transform $\mathcal{H}$: $f \mapsto \{\sin^2(\theta_p)(f)\}$, Fisher metric, generalised Euler product (M34).
- Decoherence transform $\mathcal{D}$: $f \mapsto \{S_K(f)\}$, Theorem H (monotonicity = 2nd law of sieve) (M35).

STATUS: [VAL] (32 companion scripts; see the companion scripts).

NOT PROVED: CSS quantum extension gives $d_{\text{CSS}} \sim 1$ for large n (the classical distance $d = n - 1$ does not fully transfer to the quantum setting). Cryptographic applications are speculative.

2.1 The sieve algebra

Definition 2.1 (Sieve product $*_T$)**.** For arithmetic functions $f, g\colon \mathbb{N} \to \mathbb{R}$, define:

$$(f *_T g)(n) = \sum_{d|n} f(d)\, g(n/d)\, T_3\Big(d \bmod 3,\ (n/d) \bmod 3\Big). \tag{1}$$

The sieve product weights each divisor pair by the transfer matrix entry connecting their residue classes. It inherits the divisor structure of Dirichlet convolution but modulates it by the sieve dynamics.

Sieve Algebra The algebra $(\mathcal{A}, *_T)$ of arithmetic functions under the sieve product has the following properties:

1. **Non-associative**: $(f *_T g) *_T h \neq f *_T (g *_T h)$ in general (verified by explicit counterexample on $n = 6$).
2. **Non-commutative**: $f *_T g \neq g *_T f$ in general (asymmetry of T_3: $T_{01} \neq T_{10}$).
3. **No identity**: there is no function e such that $e *_T f = f *_T e = f$ for all f.
4. **Distinct from all known algebras**: $*_T$ differs from Dirichlet convolution, Hecke algebras, and $\ell^1(\mathbb{N})$ by its non-associativity and transition-matrix weighting.

Script: `ch_math_structures/test_sieve_algebra.py` (the companion scripts).

2.2 The persistence transform

Definition 2.2 (Persistence transform P_K)**.** For an arithmetic function f and sieve level K, define $P_K(f)$ as the spectral projection of the restriction $f|_{\text{surv}(K)}$ onto the eigenbasis (v_+, v_-) of T_3:

$$P_K(f) = \langle f, v_+\rangle_K \cdot v_+ + \langle f, v_-\rangle_K \cdot v_-, \tag{2}$$

where $\langle\cdot,\cdot\rangle_K$ is the inner product weighted by the survivor distribution at level K.

Key properties:

- $P_K(\mathbf{1}) = v_+$ (the constant function is purely stationary).
- $P_K(\chi_3) = v_-$ (the Legendre character is purely anti-stationary).
- P_K is linear, idempotent ($P_K^2 = P_K$), and exact (invertible on the 2-dimensional spectral subspace).
- P_K is compatible with $*_T$: $P_K(f *_T g) = P_K(f) *_T P_K(g)$ on the eigenspace.

Script: `ch_math_structures/test_persistence_transform.py` (the companion scripts).

2.3 The PT metric and PT numbers

2.3.1 The PT metric

Definition 2.3 (PT distance)**.** For integers m, n, define:

$$d_{\mathrm{PT}}(m,n) = \sum_{K=2}^{K_{\max}} \Big|\sigma_K(m) - \sigma_K(n)\Big|, \tag{3}$$

where $\sigma_K(m) = m \bmod p_K$ is the sieve signature at level K and $K_{\max}$ is the deepest sieve level considered.

Properties of d_{PT}

1. d_{PT} is a metric (triangle inequality, symmetry, $d_{\mathrm{PT}}(m,n) = 0 \Leftrightarrow m = n$ within the primorial range).
2. d_{PT} is **quasi-orthogonal** to both the Archimedean metric $|m - n|$ and the p-adic metrics $|m - n|_p$: correlation < 0.1 in all cases.
3. Under d_{PT}, primes are **dispersed** (no two primes are close) while composites are **clustered** (sharing divisibility patterns). Classification accuracy: 96%.

Script: `ch_math_structures/test_pt_metric.py` (the companion scripts).

2.3.2 PT numbers

Definition 2.4 (PT enrichment $\mathbb{Z}_{\mathrm{PT}}$)**.** Define $\mathbb{Z}_{\mathrm{PT}}$ as the set of integers enriched by their sieve trajectory: $n \mapsto (n, \sigma_2(n), \sigma_3(n), \ldots)$. The multiplication is inherited from $\mathbb{Z}$; the addition is "disruptive" (it breaks the sieve trajectory).

Key properties: σ_K is a ring homomorphism, so both $\sigma_K(m+n) = \sigma_K(m)+\sigma_K(n) \bmod p_K$ and $\sigma_K(mn) = \sigma_K(m) \cdot \sigma_K(n) \bmod p_K$ hold exactly. However, the *sieve survivor status* (whether m is K-rough) is not preserved by addition: two K-rough integers can sum to a composite. $\mathbb{Z}_{\mathrm{PT}}$ is a *profinite filtered enrichment* of $\mathbb{Z}$.

Script: `ch_math_structures/test_pt_numbers.py` (the companion scripts).

2.4 Category of the sieve

Definition 2.5 (Cat(Sieve))**.** The *sieve category* has:

- Objects: sieve levels $K = 2, 3, 4, \ldots$
- Morphisms: transition maps $\phi_{K\to K+1}$ induced by the next prime p_{K+1}.

Four coherent functors connect the sieve to classical categories:

1. $F_{\mathrm{Grp}}\colon \mathrm{Cat}(\mathrm{Sieve}) \to \mathbf{Grp}$: maps level K to the symmetry group G_K of the transition matrix.
2. $F_{\mathrm{Vect}}\colon \mathrm{Cat}(\mathrm{Sieve}) \to \mathbf{Vect}$: maps level K to the eigenspace V_K of T_K.
3. $F_{\mathrm{Top}}\colon \mathrm{Cat}(\mathrm{Sieve}) \to \mathbf{Top}$: maps level K to the simplicial complex of the transition graph.
4. $F_{\mathrm{Info}}\colon \mathrm{Cat}(\mathrm{Sieve}) \to \mathbf{Info}$: maps level K to the divergence space $(D_{\mathrm{KL}}^{(K)}, g_F^{(K)})$.

Natural transformations between these functors encode the compatibility of algebraic, spectral, topological, and information-theoretic descriptions. The projective limit $\varprojlim_K F(K)$ converges for all four functors, giving the asymptotic sieve structure.

Script: `ch_math_structures/test_sieve_category.py` (the companion scripts).

2.5 Universal modular extension

The structures described above were constructed for the mod-3 sieve. A natural question: do they depend on $q = 3$?

Sieve Product Universality For every odd prime q, the mod-q sieve transfer matrix T_q exhibits:

1. A spectral gap: $|\lambda_2(T_q)| < 1$.
2. Forbidden self-transitions: $T_q(r, r) = 0$ for $r \neq 0$.
3. A sieve algebra $*_{T_q}$ with the same algebraic properties (non-associative, non-commutative, no identity).
4. A persistence transform $P_K^{(q)}$ with eigendecomposition.

Verified for $q = 3, 5, 7, 11, 13$. All structures are **universal**—they depend on the primality of q, not on $q = 3$ specifically. Script: `ch_math_structures/test_modular_extension.py` (the companion scripts).

2.6 PT transforms (M33–M35)

The persistence transform P_K (section 2.2) projects onto the eigenbasis of T_3, capturing the mod-3 spectral content in two coordinates per K. Three complementary transforms extend this picture by capturing *inter-modular correlations* (M33), *holonomic geometry* (M34), and *entropic monotonicity* (M35).

2.6.1 The tensor transform (M33)

Definition 2.6 (Tensor persistence transform)**.** For an arithmetic function f and sieve depth K, let $P_K^{(q)}(f) \in \mathbb{R}^q$ denote the spectral projection of f onto the eigenbasis of the transfer matrix T_q (cf. M21). The *tensor persistence transform* is

$$\mathcal{T}_K(f) = P_K^{(3)}(f) \otimes P_K^{(5)}(f) \otimes P_K^{(7)}(f) \;\in\; \mathbb{R}^{3\times5\times7}, \tag{4}$$

a third-order tensor with $3 \times 5 \times 7 = 105$ components.

The multi-modular concatenation (M30, section 2.8) stacks the three projection vectors into a 15-dimensional vector, discarding all cross-modular information. The tensor product retains the *inter-modular correlations*: how f behaves *simultaneously* mod 3, mod 5, and mod 7.

Tensor rank classification

1. $f = \mathbf{1}$: $\mathcal{T}_K(\mathbf{1})$ has rank 1 for every K (pure product, no inter-modular correlation).
2. $f = \chi_3$: rank ≤ 2 (weak coupling, mod-3 structure decouplable).
3. $f = \lambda, \mu$ (Liouville, Möbius): rank 2–3 (non-trivial arithmetic coupling between moduli).

The tensor rank measures the **complexity** of the inter-modular correlations induced by the sieve.

Plancherel identity. For a rank-1 tensor:

$$\|\mathcal{T}_K(f)\|^2 = \|P_K^{(3)}\|^2 \cdot \|P_K^{(5)}\|^2 \cdot \|P_K^{(7)}\|^2. \tag{5}$$

The coupling ratio $\rho = \|\mathcal{T}\|^2 / \prod_q \|P^{(q)}\|^2$ equals 1 exactly when the CRT factors are informationally independent; deviations quantify sieve-induced correlations.

CP decomposition. The Canonical Polyadic decomposition $\mathcal{T}_{jkl} = \sum_{r=1}^{R} a_{r,j}^{(3)} a_{r,k}^{(5)} a_{r,l}^{(7)}$ separates the tensor into R rank-1 components: mode $r = 1$ is the CRT-factorisable part; modes $r \geq 2$ are inter-modular correction terms.

Script: `ch_math_structures/test_tensor_transform.py` (the companion scripts).

2.6.2 The holonomic transform (M34)

Definition 2.7 (Holonomic transform $\mathcal{H}$)**.** For an arithmetic function f, sieve depth K, and active prime p ($p = p_j$, $j \leq K$), define the *generalised deficit*

$$\delta_p(f, K) = 1 - \frac{w_0}{\sum_{c=0}^{p-1} w_c}, \qquad w_c = \frac{n_c}{N} \langle |f|^2 \rangle_c, \tag{6}$$

where n_c counts gaps in class c mod p and $\langle |f|^2 \rangle_c$ is the conditional mean of $|f|^2$. The *holonomic angle* and its sine-squared are

$$\cos\theta_p(f) = 1 - \delta_p(f), \qquad \sin^2(\theta_p)(f) = \delta_p(f)\,(2 - \delta_p(f)). \tag{7}$$

The *holonomic transform* is the map $\mathcal{H}(f, K) = \left\{\sin^2(\theta_p)(f, K)\right\}_{p \text{ active}}$.

Properties of the holonomic transform

1. **Pythagorean identity**: $\sin^2(\theta_p) + \cos^2\theta_p = 1$ for every p (by construction).
2. **Generalised Euler product**: $\Pi(f, K) = \prod_{p \text{ active}} \sin^2(\theta_p)(f, K)$. For $f = \mathbf{1}$, this is the raw sieve holonomy. For the q_+ weights, $\Pi \to \alpha_{\text{EM}} \approx 1/137$.
3. **Fisher metric**: the infinitesimal variation

$$ds^2 = \sum_p \frac{(d\theta_p)^2}{\sin^2(\theta_p)} \tag{8}$$

is a Riemannian metric on the space of arithmetic functions, measuring informational curvature in sieve space (cf. the companion article [3]).
4. **Stability**: for fixed p, $\theta_p(f, K)$ converges when $K \gg \text{index}(p)$.
5. **Signed holonomic distance**: for functions with $|f|^2 = |g|^2$ (e.g. $f = \mathbf{1}$, $g = \chi_3$), the energy-based spectrum is identical. The *signed* distance uses class-wise means $m_c^f = \langle f \rangle_c$:

$$d_{\mathcal{H}}(f, g) = \sqrt{\sum_p \sum_{c=0}^{p-1} \frac{(m_c^f - m_c^g)^2}{p}}. \tag{9}$$

Script: `ch_math_structures/test_holonomic_transform.py` (the companion scripts).

Remark 2.8 (Circular reading of the holonomic transform). The angle θ_p is a genuine circular variable on $U(1)$, not merely a convenient parametrisation of a real quantity. This is seen from the complex lift:

$$w_p = \frac{1 - e^{2i\theta_p}}{2} = \sin^2(\theta_p)\, e^{i(\pi/2 - \theta_p)}, \tag{10}$$

which satisfies $(\operatorname{Re} w_p - \frac{1}{2})^2 + \operatorname{Im}(w_p)^2 = \frac{1}{4}$: every w_p lies on the circle of centre and radius $1/2$. The value $1/2$ appearing here is the *same* $s = 1/2$ as the stationary mod-3 occupation: both express the balance point of the sieve's forbidden-transition dynamics.

The identity $\sin^2(\theta_p) = \operatorname{Re}(1 - e^{2i\theta_p})/2$ makes the Pythagorean structure (item 1 above) a projection of circular geometry onto the real line, while the Fisher metric (8) naturally lives on the product torus $\prod_p U(1)$. A full development of the circular foundations (datum C_0: phase, winding, holonomy as primitive objects) is deferred to a companion article.

2.6.3 The decoherence transform (M35)

Definition 2.9 (Decoherence transform $\mathcal{D}$)**.** For an arithmetic function f and sieve depth K, define:

1. The *class-energy distribution*: $p_c(f, K) = \sum_{n:\,\mathrm{class}(n)=c} |f(n)|^2 \big/ \sum_n |f(n)|^2$.
2. The *quantum state* (Born rule): $|\psi_K(f)\rangle = \sum_c \sqrt{p_c}\, e^{i\phi_c}\, |c\rangle$, where $\phi_c = \arg\langle f\rangle_c$.
3. The *decoherence entropy*: $S_K(f) = -\sum_c p_c(f, K)\, \ln p_c(f, K)$.

The *decoherence transform* is $\mathcal{D}(f) = \{S_K(f)\}_{K \geq K_{\min}}$.

Decoherence Theorem (H) For every non-negative arithmetic function f and every sieve depth K:

$$S_{K+1}(f) \;\geq\; S_K(f). \tag{11}$$

This is the **second law of thermodynamics for the sieve**: the entropy of the gap-class distribution is non-decreasing under the sieve channel. Convergence is exponential, with rate controlled by the spectral gap $1 - |\lambda_2(T_3)| = 1/2 = s$ of the transfer matrix (T3).

Proof. The proof combines three ingredients already established in M1–M3: (i) the doubly stochastic structure of the sieve transfer matrix (T3), (ii) the Hardy–Littlewood–Pólya majorisation theorem, and (iii) the Schur-concavity of the Shannon entropy.

Step 1: Reduction to majorisation. Let $p^{(K)} = (p_c(f, K))_{c \in \mathbb{Z}/3\mathbb{Z}}$ denote the class-energy distribution at depth K. By construction, the application of one extra sieve step is described by the action of the doubly stochastic transfer matrix T_3 (T3, [2]) on the simplex of probability distributions on the 3 residue classes:

$$p^{(K+1)} \;=\; T_3 \cdot p^{(K)}.$$

Double stochasticity of T_3 is exact by T3 (each row *and* each column sums to 1).

Step 2: Majorisation by Hardy–Littlewood–Pólya. Let $p \succ q$ denote the standard majorisation relation: q is *majorised by* p if $\sum_{i=1}^k q_i^{\downarrow} \leq \sum_{i=1}^k p_i^{\downarrow}$ for every k, with equality at $k = n$, where $\downarrow$ denotes rearrangement in non-increasing order. The Hardy–Littlewood–Pólya theorem [1] states: for any doubly stochastic matrix M and any distribution p,

$$p \;\succ\; Mp.$$

Applied to $M = T_3$ and $p = p^{(K)}$, this gives $p^{(K)} \succ p^{(K+1)}$.

Step 3: Schur-concavity of Shannon entropy. The Shannon entropy $S(p) = -\sum_c p_c \log p_c$ is a strictly Schur-concave function on the simplex: if $p \succ q$, then $S(p) \leq S(q)$, with strict inequality unless p and q are permutations of one another. Combining with Step 2,

$$S_K(f) \;=\; S(p^{(K)}) \;\leq\; S(p^{(K+1)}) \;=\; S_{K+1}(f),$$

which is (11).

Step 4: Convergence rate. The exponential rate of approach to the equilibrium entropy $S_{\max} = \log 3$ is controlled by the second-largest eigenvalue modulus of T_3. By

T3, $\mathrm{spec}(T_3) = \{1, -1\}$, hence $|\lambda_2(T_3)| = 1$ for the bare antidiagonal. After normalisation to a row-stochastic matrix (Step 3 of the proof of T2 in [4]), the dominant non-trivial eigenvalue becomes $\lambda_2 = -1/2$, of modulus $|\lambda_2| = 1/2 = s$. Standard Perron–Frobenius theory then yields the exponential rate $\gamma_{\mathrm{dec}} = 1 - |\lambda_2| = 1/2$. □

Remark 2.10. Theorem H is unconditional: it requires no probabilistic hypothesis on f, only non-negativity. Equality $S_{K+1}(f) = S_K(f)$ holds if and only if $p^{(K)}$ is already the uniform distribution, i.e. the fixed point of T_3; in that case $S_K(f) = S_{\max} = \log 3$ and the sieve has reached informational equilibrium.

Decoherence rate. The rate $\gamma(f, K) = 1 - (S_{\max} - S_K)/(S_{\max} - S_{K-1})$ measures how fast f approaches the equilibrium entropy $S_{\max} = \ln 3$. It is tied to the spectral gap: $\gamma \sim 1 - |\lambda_2(T_3)|$.

Accessible information. The quantity $I_{\mathrm{acc}}(f, K) = S_{\max} - S_K(f) \geq 0$ is the information the sieve still preserves about f at depth K. The information loss per step $\Delta I(K) = I_{\mathrm{acc}}(K) - I_{\mathrm{acc}}(K+1) \geq 0$ is non-negative by Theorem H.

Entropic distance. The decoherence transform defines a distance combining entropy, purity, and signed class means:

$$d_{\mathcal{D}}(f, g) = \sqrt{\sum_K \Big[(S_K^f - S_K^g)^2 + (P_K^f - P_K^g)^2 + \textstyle\sum_c (m_c^f - m_c^g)^2\Big]}, \tag{12}$$

where $P_K = \mathrm{Tr}(\rho_K^2)$ is the purity and m_c the signed class mean. The signed-mean terms ensure that functions with identical $|f|^2$ (e.g. $\mathbf{1}$ and χ_3) are distinguished.

Script: `ch_math_structures/test_decoherence_transform.py` (the companion scripts).

2.6.4 Comparison of PT transforms

Property	P (M16)	$\mathcal{H}$ (M34)	$\mathcal{T}$ (M33)	$\mathcal{D}$ (M35)	Fourier
Spectral basis	$v_\pm$ of T_3	$\sin^2(\theta_p)$	$\bigotimes_q$ evecs(T_q)	entropy	e^{ikx}
Parameter	depth K	primes p	(K, q_1, q_2, q_3)	depth K	frequency k
Dim. per K	2	K (primes)	105	1	∞
Linear	YES	NO	YES	NO	YES
Monotone	NO	NO	NO	**YES**	NO
Sees sieve	**YES**	**YES**	**YES**	**YES**	NO
Inter-mod. corr.	NO	partial	**YES**	NO	NO
Curvature	NO	**YES**	NO	NO	NO

None of these transforms is a special case of Fourier or Mellin: the spectral basis is that of the sieve transfer matrix T_q, not $\{e^{ikx}\}$; the parameter is sieve depth K, not a frequency; and Theorem H has no analogue in classical harmonic analysis.

2.7 Quantum sieve and tensor networks

2.7.1 The sieve as a CPTP channel

The transfer matrix T_K can be interpreted as a quantum channel in the Kraus representation. Let ρ_K be the density matrix encoding the gap distribution at level K.

Arithmetic decoherence The sieve channel $\mathcal{E}_K : \rho \mapsto T_K \rho T_K^\dagger$ is a completely positive trace-preserving (CPTP) map with:

1. Von Neumann entropy $S_{\mathrm{vN}}(\rho_K)$ is strictly increasing with K (decoherence).
2. Purity $\mathrm{Tr}(\rho_K^2) \to 1/3$ as $K \to \infty$ (convergence to maximally mixed state on 3 classes).
3. The decoherence rate is controlled by $|\lambda_2(T_K)|$.

Script: `ch_math_structures/test_quantum_sieve.py` (the companion scripts).

This is not a claim that the sieve *is* a quantum system. It is the observation that the mathematical formalism of quantum decoherence (CPTP channels, Lindblad evolution, entropy increase) is the natural language for describing the sieve's information loss at each step.

2.7.2 CRT tensor network

The CRT factorisation $\mathbb{Z}/m\mathbb{Z} \cong \prod_{p|m} \mathbb{Z}/p\mathbb{Z}$ induces a tensor product structure on the state space. The sieve dynamics can be written as a Matrix Product State (MPS):

MPS representation The gap distribution at level K has an exact MPS representation with **bond dimension** $\chi = 3$. The entanglement entropy between any bipartition of the CRT factors satisfies $S_{\mathrm{ent}} \leq \log 3$, and the correlation length ξ is finite ($\xi < \infty$): the tensor network is **gapped**. Script: `ch_math_structures/test_tensor_network.py` (the companion scripts).

Bond dimension 3 is *exact*, not an approximation. The sieve's CRT structure is efficiently captured by a 1D tensor network with no truncation error.

2.8 Error-correcting codes from gaps

2.8.1 Multi-modular prediction

The multi-modular predictor uses residues mod $q = 3, 5, 7$ simultaneously. The mutual information between moduli is substantial (MI $\sim 1\,\mathrm{nat}$), confirming that the CRT factors are *not* informationally independent. Prediction of the next gap class improves by +61.2% over the mod-3 predictor alone.

Script: `ch_math_structures/test_multi_modular_predictor.py` (the companion scripts).

2.8.2 CRT quantum code

The CRT structure defines a quantum error-correcting code with parameters $[[3, 5.58, 1]]$. The Knill–Laflamme deviation is $\Delta_{\mathrm{KL}} = 0.008$ (approximate QEC, nearly exact). Key insight: the mutual information between moduli is *simultaneously* the resource for prediction (Direction B) and the capacity for error correction (Direction C). The two directions are the same mathematical structure viewed from different angles.

Script: `ch_math_structures/test_crt_quantum_code.py` (the companion scripts).

Honest limitation: on integer residues, the code distance is $d = 1$ (a trivial product code with no cross-modular constraints). The resolution comes from using *gap* residues.

2.8.3 The gap distance code (main theorem)

CRT-Gap Distance Theorem Let $q_1 < q_2 < \cdots < q_n$ be the first n odd primes (the moduli at sieve depth $K = n + 1$). Define the *gap residue code* $\mathcal{C}$ as:

$$\mathcal{C} = \Big\{(g \bmod q_1,\ g \bmod q_2,\ \ldots,\ g \bmod q_n) : g \in \text{distinct\_gaps}(K)\Big\}.$$

Then for $n \geq 2$:

1. The CRT map is injective on gaps: $|\mathcal{C}| = |\text{distinct\_gaps}|$.
2. The minimum Hamming distance satisfies $d = n$ for $n = 2$ (boundary case) and $d = n - 1$ for $n \geq 3$.
3. The code corrects $t = \lfloor (d-1)/2 \rfloor$ errors.
4. The ratio d/n is increasing for $n \geq 3$ and $d/n \to 1$ as $n \to \infty$ (asymptotically MDS).

Proof sketch. Suppose two distinct gaps $g_1 \neq g_2$ agree in m coordinate positions $i_1, \ldots, i_m$. Then $q_{i_1} \cdots q_{i_m}$ divides $g_1 - g_2$. Since $|g_1 - g_2| \leq \text{max\_gap} \sim 2p_K$ (linear growth) while $q_{i_1} \cdots q_{i_m} \geq 3 \cdot 5 = 15$ for $m \geq 2$ (exponential growth), we get $m \leq 1$. Hence $d \geq n - 1$. Pairs achieving exactly $n - 1$ differing positions exist (verified $K = 3 \ldots 7$), so $d = n - 1$ exactly. □

K	n (moduli)	d (distance)	t (correctable)	d/n
3	2	2	0	1.000
4	3	2	0	0.667
5	4	3	1	0.750
6	5	4	1	0.800
7	6	5	2	0.833

Comparison. A code on *integer* residues (not gap residues) always has $d = 1$: it is a product code with no cross-modular constraints. The gap residue code achieves $d = n - 1$ because gap values are bounded (linearly in p_K) while products of moduli grow exponentially. The *bound on gaps* is the ingredient that creates the distance.

Echo primes. Primes $p > p_K$ that lie below max_gap act as undetectable errors: they modify the gap sequence without changing the code coordinates. As K increases, the code distance $d = n - 1$ grows faster than the number of echo primes, so the code's error-correction capacity eventually dominates.

Script: `ch_math_structures/test_gap_distance_code.py` (the companion scripts).

2.8.4 Synthesis: four faces of the sieve

The sieve is simultaneously:

1. A prime-generating **algorithm** (number theory).
2. A CPTP **decoherence channel** (quantum information).
3. An error-correcting **code** with $d = n - 1$ (coding theory).
4. A spectral **contraction** with gap $1 - |\lambda_2|$ (dynamical systems).

The code distance $d = n - 1$ is the *arithmetic manifestation* of informational persistence: the sieve produces a code whose redundancy grows without bound, ensuring that the prime pattern is recoverable from partial information.

2.9 Benchmark and honest limits

A benchmark of the five novel mathematical structures (algebra, transform, metric, numbers, category) against concrete tasks gives:

- **Classification** (prime vs composite): d_{PT} achieves 96% accuracy, competitive with $\Omega(n)$-based classifiers.
- **Prediction**: multi-modular prediction improves over mod-3 by +61% (M30).
- **Error correction**: $d = n - 1$ proven, $t \geq 1$ for $K \geq 5$ (M32).
- **Structural detection**: the persistence transform P_K cleanly separates **1** and χ_3 (M16).

- **Categorical coherence**: all four functors converge to the asymptotic structure (M19).

Honest limits.

- d_{PT} does not replace primality tests (it classifies statistically, not deterministically).
- The gap code has $|\mathcal{C}| \sim 2n$ codewords—too few for practical communication. Its value is theoretical.
- The CPTP channel is a structural analogy, not a physical implementation.
- The CSS quantum extension gives $d_{\mathrm{CSS}} \sim 1$ for large n (the classical distance does not fully transfer).

Script: `ch_math_structures/test_capabilities.py` (the companion scripts).

Total: 32 companion scripts, all validated (the companion scripts).

3 Complex Mechanics of the Sieve

The shortest path between two truths in the real domain passes through the complex domain.
Jacques Hadamard

Status

GOAL: Lift PT from the real observable $\sin^2(\theta_p)$ to the natural complex variable $w_p = (1 - e^{2i\theta_p})/2$, uncovering hidden structure: circle geometry, holomorphic force, angular momentum, and complete Liouville integrability.

INPUTS: The companion article [3] (holonomy identity $\sin^2(\theta_p) = \delta_p(2 - \delta_p)$), The companion article [2] (GFT, action S_{PT}), the companion article [7] (α_{EM} from sieve product).

MAIN RESULTS:

- Circle theorem: all w_p lie on $\mathcal{C}(1/2, 0; s)$, radius $s = 1/2$ (Thm 3.1).
- $|w|^2 = \mathrm{Re}(w) = \sin^2(\theta_p)$: modulus squared equals real projection (ID).
- $\alpha = |W|^2 = \prod |w_p|^2$: fine structure as complex modulus (ID).
- $|T_{12,\mathbb{C}}| > T_{12,\mathrm{re}}$ by 29–62%: spectral bound strengthening.
- Holomorphic force $F(w) = i/w - 2i$ (Cauchy kernel, pole at $w = 0$).
- Angular momentum $L = \sin^2(\theta_p) = |w|^2$ is Noether charge.
- Uncertainty relation $|p_w| \cdot |w| = 1$ (exact identity, not a bound).
- Equipartition $E_p \to \kappa/2 = 1$, geometric temperature $T = 1/s = 2$.
- Complete Liouville integrability on $\mathbb{T}^n$; $\alpha = \prod L_p$.
- QED FSR coefficient: $3/(4\pi) = N_{\mathrm{spatial}} \cdot s^2/\mathrm{Circ}(\mathcal{C})$ (DER).
- Loop factor: $16\pi^2 = \Omega_2^2$ from $N_{\mathrm{spatial}} = 3$ (DER).

STATUS: [DER] + [ID] (circle theorem, 7 identities, holomorphic force, Liouville integrability, radiative corrections). 10 scripts (M41–M50), 127/127 PASS.

NOT PROVED: Physical interpretation of $\arg(W) = 0.937\pi$ (potential CP-violation link). Connection between Bernoulli–zeta corrections and existing NLO catalogue.

3.1 The complex variable w_p

At each prime p with parameter $q \in \{q_{\mathrm{stat}}, q_{\mathrm{therm}}\}$, the sieve's holonomy angle θ_p defines a natural complex variable:

$$w_p = \frac{1 - e^{2i\theta_p}}{2}. \tag{13}$$

The real and imaginary parts decompose as:

$$\mathrm{Re}(w_p) = \sin^2\theta_p = \delta_p(2-\delta_p) \qquad \text{(loss = standard PT observable)}\,, \tag{14}$$

$$\mathrm{Im}(w_p) = -\tfrac{1}{2}\sin(2\theta_p) = -\sin\theta_p\cos\theta_p \qquad \text{(coherence = hidden observable)}\,, \tag{15}$$

$$|w_p| = \sin\theta_p\,, \quad |w_p|^2 = \sin^2\theta_p = \mathrm{Re}(w_p) \qquad \text{(fundamental identity)}\,. \tag{16}$$

Standard PT uses only $\mathrm{Re}(w_p)$. The imaginary part $\mathrm{Im}(w_p) = -\sin\theta_p\cos\theta_p$ measures the *coherence* between loss and conservation channels—the geometric mean of loss and conservation amplitudes, with a sign.

Numerical values ($q = q_{\mathrm{stat}} = 13/15$, active primes):

p	$\mathrm{Re}(w_p) = \sin^2(\theta_p)$	$\mathrm{Im}(w_p)$	$\lvert w_p\rvert$
3	0.3340	−0.4716	0.5780
5	0.1992	−0.3991	0.4463
7	0.0528	−0.2237	0.2298

The coherence $|\mathrm{Im}(w_p)|$ exceeds $\mathrm{Re}(w_p)$ for $p = 5, 7$: PT standard discards the *dominant* component of the complex variable.

Circle theorem

For all primes p and all $q \in (0,1)$, the point w_p lies on the circle

$$\mathcal{C}\colon\ \left(\mathrm{Re}(w) - \tfrac{1}{2}\right)^2 + \mathrm{Im}(w)^2 = \tfrac{1}{4}\,, \tag{17}$$

centred at $(1/2, 0)$ with radius $s = 1/2$.

Proof. Let $x = \sin^2\theta$, $y = -\sin\theta\cos\theta$. Then $x^2+y^2 = \sin^4\theta+\sin^2\theta\cos^2\theta = \sin^2\theta(\sin^2\theta+\cos^2\theta) = \sin^2\theta = x$. Hence $(x-1/2)^2 + y^2 = x^2 - x + 1/4 + y^2 = 1/4$. □

The circle passes through the origin ($\theta = 0$: no loss) and through $(1,0)$ ($\theta = \pi/2$: total loss). As K increases, new primes push w_p toward the origin: *contraction is radial convergence on the circle toward $w = 0$.*

Seven fundamental identities. The following identities hold for all primes and all $q \in (0,1)$:

$$|w_p|^2 = \mathrm{Re}(w_p)\,, \tag{ID1}$$
$$\mathrm{Im}(w_p)^2 = \mathrm{Re}(w_p)\,(1 - \mathrm{Re}(w_p))\,, \tag{ID2}$$
$$w_p = -i\sin\theta_p\, e^{i\theta_p}\,, \tag{ID3}$$
$$\arg(w_p) = \theta_p - \pi/2\,, \tag{ID4}$$
$$\bar{w}_p = \frac{w_p}{2w_p - 1}\,, \tag{ID5}$$
$$z_p = 1 - 2w_p = e^{2i\theta_p}\,, \tag{ID6}$$
$$\delta_p = \frac{(1-q)[p]_q}{p}\,. \tag{ID7}$$

Identity ID1 is the defining property of the circle $\mathcal{C}$. Identity ID5 makes the conjugate a *Möbius transform*, rendering the force $F = -i/\bar{w}$ holomorphic in w alone (Section 3.5).

3.2 Complex observables

3.2.1 Complex product and α

The sieve's fine-structure constant arises from the product:

$$W \;=\; \prod_{p\text{ active}} w_p\,, \qquad |W|^2 = \prod_p |w_p|^2 = \prod_p \sin^2(\theta_p) = \alpha_{\mathrm{bare}}\,. \tag{18}$$

The *phase* of W is a new observable, independent of $|W|^2 = \alpha$:

$$\arg(W) = \sum_p \arg(w_p) = \sum_p (\theta_p - \pi/2)\,. \tag{19}$$

For $q = q_{\text{stat}}$, $\arg(W) \approx 0.937\pi$. This is *not* a simple fraction of π; the closest rational approximation is 15/16.

3.2.2 Complex T_{12} and spectral strengthening

The spectral gap T_{12} admits a complex extension via the character χ_3:

$$T_{12,\mathbb{C}} = \frac{1}{3}\sum_p \chi_3(p)\, w_p = \frac{w_7 - w_5}{3}\,, \tag{20}$$

since $\chi_3(3) = 0$, $\chi_3(5) = -1$, $\chi_3(7) = +1$. The modulus satisfies:

$$|T_{12,\mathbb{C}}| > |T_{12,\text{re}}|\,, \qquad \text{ratio } \frac{|T_{12,\mathbb{C}}|}{|T_{12,\text{re}}|} = \frac{1}{\sin\theta_{\text{eff}}} \in [1.29,\ 1.62]\,.$$

For q_{stat}: $|T_{12,\mathbb{C}}| = 0.00920$ vs $T_{12,\text{re}} = 0.00712$ (+29%). For q_{therm}: gain +62%. This strengthens spectral bounds in the T5 convergence proof (the companion article [4]) without changing the proof structure.

3.2.3 Complex action

The PT effective action $S_{\text{PT}} = -\sum_p \ln\sin^2(\theta_p)$ extends to:

$$S_{\mathbb{C}} = -\sum_p \ln w_p = -\sum_p \Big[\ln|w_p| + i\arg(w_p)\Big]\,, \tag{21}$$

with $\text{Re}(S_{\mathbb{C}}) = S_{\text{PT}}/2$ (since $\ln\sin^2(\theta_p) = 2\ln|w_p|$). The imaginary part encodes the accumulated phase along the path of primes.

3.3 Deep algebraic structure

3.3.1 U(1) correspondence

The affine map $z_p = 1 - 2w_p = e^{2i\theta_p}$ sends the circle $\mathcal{C}$ to the unit circle $\mathbb{S}^1$. This identifies the sieve state space with U(1).

3.3.2 Fisher = Fubini–Study

The Fisher metric on the Bernoulli parameter space pulls back to four times the arc-length metric on $\mathcal{C}$:

$$\mathrm{d}s^2_{\text{Fisher}} = \frac{4}{\sin^2(\theta_p)(1-\sin^2(\theta_p))}\,(\mathrm{d}\sin^2(\theta_p))^2 = 4\,\mathrm{d}\theta_p^2\,. \tag{22}$$

This is the Fubini–Study metric on $\mathbb{CP}^1$. The factor 4 equals the inverse of the circle's radius squared: $4 = 1/s^2$.

3.3.3 Density matrix correspondence

Each w_p encodes the first column of a rank-1 density matrix $\rho_p = |\psi_p\rangle\langle\psi_p|$ for a 2-level system $|\psi_p\rangle = \sin\theta_p\,|0\rangle + \cos\theta_p\,|1\rangle$:

$$w_p = (\rho_p)_{00} - i\,(\rho_p)_{01} = \sin^2\theta_p - i\sin\theta_p\cos\theta_p\,. \tag{23}$$

The real part is the diagonal element (probability), the imaginary part is the off-diagonal element (coherence). The circle $\mathcal{C}$ is the locus of pure states on the Bloch sphere's meridian $\Phi = 0$.

3.3.4 Cross-ratios

For four co-cyclic points on $\mathcal{C}$, the cross-ratio is *exactly real* (imaginary part $< 10^{-16}$). For the active primes:

$$\mathrm{CR}(3, 5;\, 7, \infty) \approx 2.000 \quad (\text{error } 1.5 \times 10^{-4})\,.$$

The quasi-integer cross-ratios reflect the harmonic structure of the first three active primes.

3.4 Complex corrections

3.4.1 Circle radius encodes s

The circle $\mathcal{C}$ has radius $s = 1/2$. The fundamental parameter of PT is thus encoded in the geometry of the complex plane: *the geometry is the theory.*

3.4.2 Structural decomposition of $1/\alpha$

The bare product $1/\alpha_{\text{bare}} = \prod(1/\sin^2(\theta_p))$ decomposes via $\sin^2 = \theta^2 \cdot (\sin\theta/\theta)^2$:

$$\frac{1}{\alpha_{\text{bare}}} = \underbrace{\prod_p \frac{1}{\theta_p^2}}_{=110.34 \text{ (classical)}} \times \underbrace{\prod_p \left(\frac{\theta_p}{\sin\theta_p}\right)^2}_{=1.235 \text{ (Bernoulli–zeta)}}\,. \tag{24}$$

The "classical" factor uses only the leading-order $\sin\theta \approx \theta$. The Bernoulli–zeta factor encodes higher-order corrections via $\theta/\sin\theta = 1 + \theta^2/6 + 7\theta^4/360 + \cdots$, where the coefficients involve Bernoulli numbers B_{2k} and implicitly $\zeta(2k)$.

This decomposition is an identity—the Bernoulli factor is already contained in $\sin^2$ and does not constitute a new correction to α. It reveals the internal structure of the bare product.

3.4.3 NLO from coherence

The ratio of imaginary to real parts gives the NLO/LO ratio:

$$\frac{|\mathrm{Im}(w_p)|}{\mathrm{Re}(w_p)} = |\cot\theta_p| \sim \sqrt{p/2}\,. \tag{25}$$

For all primes beyond $p = 3$, the coherence (imaginary) component exceeds the loss (real) component: NLO *dominates* LO in the complex plane.

3.4.4 Radiative corrections from circle geometry

The last two QFT formulas used in the NLO catalogue (the companion article [7], Appendix B)—the QED final-state radiation coefficient and the one-loop vertex factor—decompose entirely into PT-native quantities.

[DER]

QED FSR from circle geometry The QED final-state radiation correction for a fermion of charge Q_f is:

$$\delta_{\mathrm{QED}} = \frac{3\alpha Q_f^2}{4\pi} = \frac{N_{\mathrm{spatial}} \cdot s^2}{\mathrm{Circ}(\mathcal{C})} \cdot \alpha\, Q_f^2\,, \tag{26}$$

where $N_{\mathrm{spatial}} = 3$ (Theorem, R38b), $s = 1/2$ (fundamental), and $\mathrm{Circ}(\mathcal{C}) = 2\pi s = \pi$. The coefficient $3/4 = N_{\mathrm{spatial}} \cdot s^2$ is thus determined by the number of spatial dimensions and the sieve radius.

Physical interpretation: N_{spatial} counts the independent polarisation directions for the emitted quantum; $s^2 = 1/4$ is the ratio $\mathrm{Area}(\mathcal{C})/\pi$; and $\mathrm{Circ}(\mathcal{C}) = \pi$ normalises the emission phase space on the circle.

[DER]

Loop factor from solid angle The one-loop normalisation factor is:

$$16\pi^2 = \Omega_2^2 = \left(\frac{2\pi^{N_{\mathrm{spatial}}/2}}{\Gamma(N_{\mathrm{spatial}}/2)}\right)^2, \tag{27}$$

where $\Omega_2 = 4\pi$ is the solid angle of the unit S^2 in $N_{\mathrm{spatial}} = 3$ spatial dimensions.

The vertex correction $\rho_b = 1 - G_F m_t^2/(4\sqrt{2}\,\pi^2)$ (R43, non-universal $Z \to b\bar{b}$) rewrites in Yukawa form:

$$\rho_b = 1 - \frac{y_t^2}{\Omega_2^2}\,, \qquad y_t = \frac{\sqrt{2}\,m_t}{v}\,, \tag{28}$$

where every quantity is PT-derived: y_t from the quark mass hierarchy, v (R51), and $\Omega_2^2 = 16\pi^2$ from $N_{\mathrm{spatial}} = 3$. The numerical value $\rho_b = 0.9938$ matches the standard result exactly (algebraic identity).

Contour integral. The holomorphic force $F(w) = i/w - 2i$ satisfies:

$$\oint_{\mathcal{C}} F \,\mathrm{d}w = -\pi = -\mathrm{Circ}(\mathcal{C})\,. \tag{29}$$

Since $w = 0$ lies *on* $\mathcal{C}$ (boundary pole), the Sokhotski–Plemelj formula gives $\oint(i/w)\,\mathrm{d}w = \pi i \cdot i = -\pi$ (half the full residue). The total work done by the force along one traversal of $\mathcal{C}$ equals its circumference: a topological identity linking radiation to geometry.

Status. With these two derivations, *all NLO/NNLO corrections used in the companion article* [7] *are now PT-derived*; no external QFT formulas remain in the dressing chain. (Script: M50, `test_complex_radiative`, 18/18 PASS.)

3.5 Holomorphic mechanics

3.5.1 The force

Define the complex momentum as $p_w = \mathrm{d}S_{\mathbb{C}}/\mathrm{d}\theta = -\cot\theta - i$. The associated force in w-space is:

$$F(w) = \frac{i}{w} - 2i\,. \tag{30}$$

This is a holomorphic function on $\mathbb{C} \setminus \{0\}$ with a simple pole at $w = 0$ and residue $\mathrm{Res}(F, 0) = i$. The identity ID5 $\bar{w} = w/(2w - 1)$ converts $F = -i/\bar{w}$ into the holomorphic form above.

Rotation identity.

$$F(w_p) \cdot \mathrm{Re}(w_p) = -i\, w_p\,. \tag{31}$$

The force times the observable equals the complex variable rotated by $-\pi/2$. Multiplication by $\sin^2(\theta_p)$ performs a quarter-turn rotation.

Universal azimuthal force.

$$\mathrm{Im}\left(\frac{\mathrm{d}S_{\mathbb{C}}}{\mathrm{d}\theta}\right) = -1 = -\frac{1}{2s}\,, \tag{32}$$

constant for all primes. The azimuthal force equals the inverse of the circle diameter $2s = 1$.

3.5.2 Angular momentum as Noether charge

The angular momentum associated to the U(1) rotation $w \mapsto e^{i\alpha} w$ is:

$$L_p = \mathrm{Im}(\bar{w}_p \, \mathrm{d}w_p/\mathrm{d}\theta) = \sin^2(\theta_p) = |w_p|^2 \,. \tag{33}$$

The PT observable IS angular momentum.

3.5.3 Uncertainty relation

$$|p_w| \cdot |w_p| = 1 \qquad \text{for all } p \,. \tag{34}$$

This is an *exact identity* (not a bound), with constant equal to the diameter $2s = 1$ of the circle $\mathcal{C}$. From (34): $\sin^2(\theta_p) = |w|^2 = 1/(1 + \cot^2 \theta) = 1/(1 + \mathrm{Re}(p_w)^2)$, recovering the Born rule from momentum.

3.6 Integrability and action–angle variables

3.6.1 Liouville actions

Define the action variable and Hamiltonian:

$$L_p = \sin^2(\theta_p) \,, \qquad H_p = -\ln \sin \theta_p = -\tfrac{1}{2} \ln L_p \,. \tag{35}$$

The Poisson bracket $\{L_p, H_{p'}\} = 0$ for $p \neq p'$ (independence of primes), and $\{L_p, H_p\} = 0$ since both depend only on θ_p (one degree of freedom per prime). The system is *completely integrable* in the sense of Liouville, living on a torus $\mathbb{T}^n$.

The fine-structure constant is the product of Liouville actions:

$$\alpha = \prod_{p \text{ active}} L_p = \prod_{p \text{ active}} \sin^2(\theta_p) \,. \tag{36}$$

3.6.2 Equipartition

Define the kinetic energy $E_p = \frac{1}{2} p \, \theta_p^2$. For large p, $\theta_p \sim \sqrt{2/p}$, so:

$$E_p \;\to\; \frac{\kappa}{2} = 1 \,, \qquad T_{\text{geom}} = \kappa = \frac{1}{s} = 2 \,. \tag{37}$$

Exact correction: $E_p = 1 - 2/p + O(1/p^2)$. The geometric temperature $T = 2$ equals the curvature of the circle $\mathcal{C}$ (inverse of radius $s = 1/2$).

3.6.3 Bernoulli–zeta corrections

The ratio $\theta/\sin\theta$ admits the Bernoulli series:

$$\frac{\theta}{\sin\theta} = 1 + \frac{\zeta(2)}{\pi^2}\,\theta^2 + \frac{7\zeta(4)}{2\pi^4}\,\theta^4 + \cdots = \sum_{k=0}^{\infty} \frac{(-1)^{k+1}(2^{2k}-2)\,B_{2k}}{(2k)!}\,\theta^{2k}\,. \tag{38}$$

This connects the "quantum correction" factor $\prod(\theta/\sin\theta)^2 = 1.235$ in (24) to the values of the Riemann zeta function at even integers.

3.6.4 UV regularisation

The discreteness of primes provides a natural UV cutoff: $p_{\min} = 3$ (the smallest active prime). The echo primes ($p = 11, 13$) serve as counter-terms in the effective action, analogous to Pauli–Villars regulators. There are *no divergences, no poles, and no cutoff dependence*: the complex mechanics is finite by construction.

3.7 Synthesis

3.7.1 Real PT vs. Complex PT

Concept	**Real PT**	**Complex PT**
Observable	$\sin^2(\theta_p) \in [0,1]$	$w_p \in \mathcal{C} \subset \mathbb{C}$
Geometry	Interval $[0,1]$	Circle of radius $s = 1/2$
Metric	Fisher $4\,\mathrm{d}\theta^2$	Fubini–Study on $\mathbb{CP}^1$
Coherence	Discarded	$\mathrm{Im}(w_p) = -\sin\theta\cos\theta$
α	$\prod\sin^2(\theta_p)$	$\lvert W\rvert^2$, plus phase $\arg(W)$
Action	$S = -\sum\ln\sin^2(\theta_p)$	$S_{\mathbb{C}} = -\sum\ln w_p$
Force	$-\mathrm{d}S/\mathrm{d}\theta$ (real)	$F(w) = i/w - 2i$ (holomorphic)
Angular mom.	Not defined	$L = \sin^2(\theta_p)$ (Noether charge)
Uncertainty	Not defined	$\lvert p\rvert \cdot \lvert w\rvert = 1$ (identity)
Integrability	Not manifest	Liouville on $\mathbb{T}^n$
Temperature	$\beta = 1/\mu$	$T = 1/s = 2$ (curvature)

3.7.2 Impact on existing proofs

1. **Spectral bounds** (the companion article [4]): $|T_{12,\mathbb{C}}| > T_{12,\mathrm{re}}$ tightens the spectral gap by 29–62%. The T4/T5 proofs remain valid and are strengthened.
2. **Fine structure** (the companion article [7]): $\alpha = |W|^2$ provides a geometric interpretation; the structural decomposition (24) reveals the Bernoulli–zeta content already present in $\sin^2(\theta_p)$. No numerical change.

3. **Fisher metric** (the companion article [3]): The metric $4\,\mathrm{d}\theta^2$ is Fubini–Study on the circle. This is a reformulation, not a correction.
4. **Born rule** (the companion article [7]): $\sin^2(\theta_p) = |w|^2$ is literally $|\psi|^2$, and the density matrix correspondence (23) gives an explicit quantum embedding.

3.7.3 Open directions

1. Physical meaning of $\arg(W) = 0.937\pi$: link to Jarlskog invariant or CP violation?
2. Berry phase accumulation as K increases: quantisation of $\gamma_K = \sum \mathrm{Im}\langle\psi_k|\psi_{k+1}\rangle$?
3. Further radiative processes derivable from the contour integral $\oint_{\mathcal{C}} F\,\mathrm{d}w = -\pi$ and the Cauchy kernel structure of F.
4. q-deformation and circle breaking: when $q \to 1$, the circle collapses to a point; when $q \to 0$, it expands to the full disk.

Scripts. M41 (`test_imaginary_PT`), M42 (`test_complex_plane`), M43 (`test_complex_deep`), M44 (`test_complex_directions`), M45 (`test_complex_corrections`), M46 (`test_complex_questio`
M47 (`test_complex_force`), M48 (`test_complex_mechanics`), M49 (`test_complex_final`), M50 (`test_complex_radiative`). All in `scripts/ch_math_structures/`.

4 Predictive Mathematics

> Give me a field and a constraint cascade. I will tell you how many degrees of freedom survive at each level, what their activation cost is, and what geometry they carry.
>
> The PM programme

STATUS

GOAL: Develop the *Predictive Mathematics* (PM) framework: a universal language for predicting how constraint cascades shape the degrees of freedom that survive. Instantiate the framework on the Eratosthenes sieve and derive closed-form theorems for layer dimension, activation spectrum, and activation threshold, including the phantom rank mechanism.

INPUTS: Holonomy angles $\sin^2(\theta_p)$ and anomalous dimensions γ_p (the companion article [3]), fixed point $\mu^* = 15$ (the companion article [4]), $N_{\text{spatial}} = 3$ (the companion article [4]), CRT factorisation (the companion article [2]).

MAIN RESULT:

- L1: each prime adds exactly 1 DOF (Theorem 4.3, unconditional [THM]).
- $\dim(d) = P(d{+}1) - (2d{+}1)$, layer dimension formula (Theorem 4.4, [THM] under gen. pos.).
- $\Sigma_k = \binom{4}{k}$ for $k \geq \text{AT}$, activation spectrum (Theorem 4.5, [THM] under gen. pos.).
- AT formula in three regimes: echo, construction, asymptotic (Theorems 4.6.1–4.6.3, [THM]).
- Partition identity: $\text{phantom}(1) + \text{struct}(1) = n_{\text{base}}$ (Identity 4.6.3, [ID]).
- PM–PT conjecture proved: $\text{AT} = 1$ for all $p \geq 41$ (Corollary 4.12).

STATUS: [THM] (L1, AT parts A–B) + [THM] (dim, Σ, AT part C, under gen. pos.) + [ID] (C1)

NOT PROVED: The general position conjecture (Conjecture 4.8) is verified numerically for all 13 shells (11–59) but not derived from arithmetic first principles. The second instantiation (error-correcting codes) is [VAL], not [THM].

4.1 The PM framework

4.1.1 Motivation

The Theory of Persistence, as developed in The companion articles [2–4], proves that the Eratosthenes sieve converges to a unique fixed point $\mu^* = 15$ with active primes $\{3, 5, 7\}$.

The bridge axioms (the companion article [5]) then map this structure to physics. But one question remains untouched:

> *What happens, structurally, when a prime is added to the sieve base?*

The *Predictive Mathematics* (PM) framework answers this question. It is not a mathematics of objects—it predicts what *becomes inevitable* when constraints are stacked on a raw field. Three outputs:

1. the number of surviving degrees of freedom (DOF) at each constraint level,
2. the activation cost to unlock the next DOF,
3. the geometry carried by the survivors (transversality spectrum, Fisher metric, persistence invariants).

4.1.2 Universal definitions

Definition 4.1 (Constraint cascade)**.** A *constraint cascade* is a tuple (F, C, S, G, I) where F is a raw field of candidates, $C = (C_1, C_2, \ldots)$ is an ordered family of constraints, $S_d = \{x \in F : C_1(x) \wedge \cdots \wedge C_d(x)\}$ is the survivor set at depth d, G is the geometry readable from the survivor distribution, and I is a family of invariants stable under the cascade.

Definition 4.2 (Observation matrix)**.** Given a family of probes $\eta = (\eta_1, \ldots, \eta_n)$ and a depth d, the *observation matrix* $\mathbf{O}(\eta, d)$ is the centred Gram matrix whose rows are the conditional expectations $\mathbb{E}[\eta_j \mid \text{conditioning level } d]$ for each conditioning residue. The *historical matrix* at depth d is the vertical stack

$$\mathbf{H}(\eta, d) = \begin{pmatrix} \mathbf{O}(\eta, 1) \\ \vdots \\ \mathbf{O}(\eta, d) \end{pmatrix}.$$

The *cumulative rank* is $R(\eta, d) = \operatorname{rank} \mathbf{H}(\eta, d)$.

Definition 4.3 (Layer dimension and activation threshold)**.** The *layer dimension* at depth d is $\dim(d) = R(\eta, d) - R(\eta, d{-}1)$. The *future depth* of a shell p is the largest d with $\dim(d) > 0$. The *activation threshold* AT is the minimum number of target probes required to detect the new DOF at the future depth.

Definition 4.4 (Transversality spectrum)**.** For a probe packet of size n_p, the *transversality spectrum* $(\Sigma_1, \ldots, \Sigma_{n_p})$ counts how many subsets of size k detect the new DOF: $\Sigma_k = \#\{S \subseteq \text{packet} : |S| = k,\ S \text{ detects}\}$.

4.2 Sieve instantiation

4.2.1 Setup

In the sieve instantiation:

- $F = \{n \in \mathbb{N} : n \geq 2\}$ (candidate integers).
- C_k: divisibility by the k-th prime p_k (sieve of Eratosthenes).
- S_d: the d-rough integers (coprime to $p_1, \ldots, p_d$).
- Probes: centred indicators $\eta(p,t)(g) = \delta(g \equiv t \bmod p) - 1/p$ for each prime p and residue $t \in \{p-4, \ldots, p-1\}$ (a packet of 4 probes per shell).
- Conditioning: residues modulo the active primes $\{3, 5, 7\}$ and echo primes $\{11, 13, \ldots\}$ already in the base.

Each prime $p > 7$ added to the sieve base defines a *shell*. The shell explores how the PM invariants—future depth, activation threshold, transversality spectrum—change when the constraint C_p is imposed.

4.2.2 Primorial clock

The shells tile the primorial cycle $\mathbb{Z}/30\mathbb{Z}$, where $30 = 2\mu^*$ is the primorial of the active primes. The 8 coprime residues $(\mathbb{Z}/30\mathbb{Z})^* = \{1, 7, 11, 13, 17, 19, 23, 29\}$ are visited in order by the primes > 7. The first complete cycle (shells 11–37) covers all 8 residues; cycle 2 begins at shell 41.

4.2.3 Data: 13 shells explored

Table 1: PM invariants for all 13 explored shells.

Shell	Depth	Incr.	AT	Spectrum $(\Sigma_1, \ldots, \Sigma_4)$	Cycle
11	4	1	1	(4, 6, 4, 1)	1
13	5	1	1	(4, 6, 4, 1)	1
17	6	1	2	(0, 6, 4, 1)	1
19	7	1	2	(0, 6, 4, 1)	1
23	8	1	2	(0, 6, 4, 1)	1
29	9	1	3	(0, 0, 4, 1)	1
31	10	1	3	(0, 0, 4, 1)	1
37	11	1	3	(0, 0, 4, 1)	1
41	12	1	1	(4, 6, 4, 1)	2
43	13	1	1	(4, 6, 4, 1)	2
47	14	1	1	(4, 6, 4, 1)	2
53	15	1	1	(4, 6, 4, 1)	2
59	16	1	1	(4, 6, 4, 1)	2

Three facts are visible:

1. The increment is always 1 (L1).

2. The spectrum is always $\binom{4}{k}$ for $k \geq \mathrm{AT}$, zero below.
3. The AT follows an *escalator* pattern in cycle 1 ($\mathrm{AT} = 1, 1, 2, 2, 2, 3, 3, 3$) and resets to $\mathrm{AT} = 1$ for all of cycle 2.

4.3 Theorem L1: unit increment

L1 — Unit Increment Theorem For every prime $p > 7$, adding p to the sieve base increases the future depth by exactly 1:

$$\mathrm{FD}(\text{shell } p) = \text{base\_depth} + 1.$$

The proof rests on two algebraic lemmas.

Lemma 4.5 (Indicator saturation)**.** *For any prime p and distinct residues $t_1 \neq t_2 \in \mathbb{Z}/p\mathbb{Z}$,*

$$\delta(g \equiv t_1) \cdot \delta(g \equiv t_2) = 0 \qquad \textit{(identically, not statistically).}$$

Moreover $\delta(g \equiv t)^2 = \delta(g \equiv t)$ (idempotence).

Proof. A gap g is congruent to exactly one residue modulo p, so the product of indicators for distinct residues vanishes identically. Idempotence follows from $0^2 = 0$, $1^2 = 1$. □

Consequence. Every monomial of degree ≥ 2 in the probes of a *single* prime is a linear combination of degree-1 probes (plus constants). The probe algebra for prime p is spanned by the $p-1$ centred indicators $\{\eta(p,t)\}_{t \neq 0}$.

Lemma 4.6 (CRT irreducibility)**.** *For distinct primes $p \neq q$, the product $\delta(g \equiv t \bmod p) \cdot \delta(g \equiv s \bmod q)$ equals $\delta(g \equiv r \bmod pq)$ where r is the unique CRT lift. This observable at module pq is* not *in the span of observables at modules p and q separately.*

Proof. Suppose $\delta(g \equiv r \bmod pq) = \sum_i a_i f_i(g \bmod p) + \sum_j b_j h_j(g \bmod q) + c$. Evaluating at $g = r$ and $g = r + p$ (same residue mod p, different mod q) forces $b_j = 0$ for all j. By symmetry, $a_i = 0$ for all i. But $\delta(g \equiv r \bmod pq)$ is non-constant: contradiction. □

Proof of Theorem 4.3. **Part I:** $\mathrm{FD} \geq \text{base\_depth} + 1$**.** At depth $D = d{+}1$ (after adding $p = p_{d+1}$), the observables include cross-products $\eta(p_1, t_1) \cdots \eta(p_d, t_d) \cdot \eta(p, t)$—a product of $d{+}1$ probes for *distinct* primes. By iterated CRT ($\mathbb{Z}/m'\mathbb{Z} \cong \prod \mathbb{Z}/p_i\mathbb{Z}$), this product contains the irreducible observable $\delta(g \equiv r \bmod m')$ at module $m' = p_1 \cdots p_d \cdot p$. Since $D_{\mathrm{KL}}(P_p \| U_p) > 0$ (the bias of gaps mod p is non-trivial by T0), the $\mathbb{Z}/p\mathbb{Z}$ component carries non-zero information. The rank at depth $d{+}1$ is strictly larger than at depth d.

Part II: $\mathrm{FD} \leq \text{base\_depth} + 1$**.** At depth $D = d{+}2$, every monomial of degree $d{+}2$ in the probes of $d{+}1$ distinct primes must involve at least one prime squared (pigeonhole). By Lemma 4.5, the squared factor reduces to degree ≤ 1, collapsing the monomial to degree $\leq d{+}1$. No new direction appears at depth $d{+}2$. □

Remark 4.7. L1 is unconditional: no probabilistic hypothesis, no general position assumption. It holds for *any* sequence of gaps, not only prime gaps. The theorem is a property of the CRT algebra of modular indicators.

4.4 Layer dimension formula

Conjecture 4.8 (General position). The interaction corrections between distinct primes in the observation matrix are in general position: no accidental linear relations arise among the CRT cross-terms. (Verified for 13 shells, not proved.)

Dimension Formula Under Conjecture 4.8, the stabilised dimension of the layer at depth d is

$$\dim(d) = P(d+1) - (2d+1),$$

where $P(n) = \sum_{k=1}^{n} p_k$ is the sum of the first n primes. Equivalently,

$$\dim(d) = \sum_{k=1}^{d} (p'_k - 1) - (d - 1),$$

where $p'_k = p_{k+1}$ is the k-th odd prime.

Proof sketch. The observation matrix at depth d has d blocks of conditioning rows, one per odd prime $p'_k \in \{3, 5, \ldots, p'_d\}$. After centring, block k contributes $p'_k - 1$ effective rows. The total number of effective rows is $\sum_{k=1}^{d}(p'_k - 1)$.

Among these rows, exactly $d-1$ inter-block dependencies exist: each of the $d-1$ "old" blocks (those already present at depth $d-1$) shares one marginal direction with the previous-depth space V_{d-1}. The "new" block (for prime p'_d, entering at depth d) introduces no such dependency. Under general position, these are the *only* dependencies. Hence $\dim(d) = \sum(p'_k - 1) - (d-1)$. □

Increment. An immediate consequence is $\Delta \dim = \dim(d+1) - \dim(d) = p_{d+2} - 2$, the number of non-trivial transitions in the transfer matrix $T_{p_{d+2}}$ of size $(p_{d+2}-1) \times (p_{d+2}-1)$. Table 2 verifies the dimension formula for $d = 2, \ldots, 11$.

Look-ahead. The offset of 2 (the layer at depth d depends on prime p_{d+2}) coincides with DEPTH $= 2$, the sieve depth from Theorem T4.

Prediction: $\dim(12) = P(13) - 25 = 238 - 25 = 213$.

Table 2: Layer dimension verification (10/10).

d	$P(d+1)$	$2d+1$	$\dim(d)$ predicted	observed	match
2	10	5	5	5	✓
3	17	7	10	10	✓
4	28	9	19	19	✓
5	41	11	30	30	✓
6	58	13	45	45	✓
7	77	15	62	62	✓
8	100	17	83	83	✓
9	129	19	110	110	✓
10	160	21	139	139	✓
11	197	23	174	174	✓

4.5 Activation spectrum

Transversality Spectrum Under Conjecture 4.8 (AT-general position), for a probe packet of size 4:

$$\Sigma_k = \begin{cases} 0 & \text{if } k < \mathrm{AT}, \\ \binom{4}{k} & \text{if } k \geq \mathrm{AT}. \end{cases}$$

Proof. Let $W = \operatorname{span}(u_1, \ldots, u_4)$ be the transverse space (the projections of the 4 probes onto $V^\perp$), with $\dim W = \mathrm{AT}$.

Case $k < \mathrm{AT}$: A subset of $k < \mathrm{AT}$ probes spans a space of dimension $\leq k < \mathrm{AT}$ and cannot cover W.

Case $k \geq \mathrm{AT}$: Every subset of size k contains a sub-subset of size AT which, by general position, has rank exactly AT and spans W. Hence every such subset detects the DOF, giving $\Sigma_k = \binom{4}{k}$. □

Verification: 13/13 exact (Table 1).

4.6 Activation threshold: three regimes

The activation threshold is governed by $|B| = |B(p)| = \pi(p) - 4$, the number of echo conditioning levels (primes between 7 and p, exclusive).

4.6.1 Part A: CRT perturbative reset (cycles ≥ 2)

Activation Threshold A If $|B(p)| > \varphi(30) = 8$ (equivalently: there exists $q < p$ with $q \equiv p \pmod{30}$), then $\mathrm{AT}(p) = 1$.

Proof. For $|B| \geq 9$ (i.e. $p \geq 41$), there exists a predecessor $q < p$ with $q \equiv p \pmod{30}$, i.e. $q \equiv p \pmod 2$, $q \equiv p \pmod 3$, $q \equiv p \pmod 5$. The CRT component mod 30 of $\eta_{p,c}$ is already captured by the existing probe $\eta_{q,c'}$ in the base. The complementary component (distinguishing p from q within their shared mod-30 channel) is rank-1. A single probe $\eta_{p,c}$ suffices to capture this rank-1 correction. By L1, the future depth adds exactly one DOF, so $\mathrm{AT} = 1$. □

4.6.2 Part B: echo regime ($|B| \leq 2$)

Activation Threshold B For $p \in \{11, 13\}$ (echo primes, $|B| \leq 2$): $\mathrm{AT}(p) = 1$.

Proof. The primes 11 and 13 are the first echo primes ($\gamma_{11} = 0.426$, $\gamma_{13} = 0.356$, both $< 1/2$). At their depth, the base contains no probe for residues mod 11 (resp. 13). The entire probe space $\eta_{p,c}$ is new. A single probe captures the unique new DOF guaranteed by L1. □

4.6.3 Part C: phantom rank mechanism (cycle 1)

The heart of the chapter. We first establish the *phantom rank* phenomenon.

Definition 4.9 (Phantom and structural rank)**.** Let F_0 be the base family (n_{base} probes) and $F_k = F_0 \cup \{\eta_1, \ldots, \eta_k\}$ the extended family with k target probes. Define:

$$\mathrm{phantom}(k) = R(F_k, d) - R(F_0, d), \tag{39}$$

$$\mathrm{struct}(k) = R(F_k, d{+}1) - R(F_k, d), \tag{40}$$

where $d = \text{base\_depth}$ and $d{+}1$ is the future depth. Then $\mathrm{AT} = \min\{k : \mathrm{struct}(k) > 0\}$.

[ID]

C1 — Partition Identity $\mathrm{phantom}(1) + \mathrm{struct}(1) = n_{\text{base}}$.

Proof. The base family observes all depths up to d and contains no probe for the target prime p. Hence the base-only structural rank is zero: $\dim_{\text{base}} = R(F_0, d{+}1) - R(F_0, d) = 0$. Then:

$$\begin{aligned}
\mathrm{phantom}(1) + \mathrm{struct}(1) &= [R(F_1, d) - R(F_0, d)] + [R(F_1, d{+}1) - R(F_1, d)] \\
&= R(F_1, d{+}1) - R(F_0, d) \\
&= [R(F_1, d{+}1) - R(F_0, d{+}1)] + \dim_{\text{base}} \\
&= R(F_1, d{+}1) - R(F_0, d{+}1).
\end{aligned}$$

Under general position, the n_{base} cross-terms ($\eta_{\text{target}} \times \text{base}_j$) at depth $d{+}1$ are linearly independent, so $R(F_1, d{+}1) - R(F_0, d{+}1) = n_{\text{base}}$. □

Table 3 confirms the partition identity for all six shells tested.

Table 3: Partition identity C1: exact verification (6/6).

Shell	n_{base}	phantom(1)	struct(1)	ph + st	match
11	6	4	2	6	✓
13	10	8	2	10	✓
17	14	14	0	14	✓
19	18	18	0	18	✓
23	22	22	0	22	✓
29	26	26	0	26	✓

C2 — Phantom Saturation Under Conjecture 4.8, for $|B| \geq 3$: $\text{phantom}(1) = n_{\text{base}}$, hence $\text{struct}(1) = 0$.

Proof. When p does not divide $M_d = \prod_{k=1}^{d} p_k$, the distribution of $\text{surv}(d) \bmod p$ is non-uniform. The CRT decomposition of the observation matrix reveals $N_{\text{spatial}} = 3$ active channels (mod 3, mod 5, mod 7). The non-uniformity creates correlations between $\eta_{p,c}$ and base probes at depth d.

For $|B| \leq 2$: only 1–2 echo conditioning levels; the 3 CRT channels are not all saturated. Two structural directions survive: $\text{struct}(1) = 2 > 0$.

For $|B| \geq 3$: the $|B|$ echo levels provide enough non-uniformity to saturate all 3 CRT channels. All n_{base} cross-terms are absorbed by the phantom: $\text{struct}(1) = 0$.

The threshold $|B| = 3$ is a dimensional count: it takes at least $N_{\text{spatial}} = 3$ independent sources of non-uniformity to block all 3 CRT channels. □

Activation Threshold C Under Conjecture 4.8, for non-echo shells of cycle 1 ($3 \leq |B| \leq 8$):

$$\text{AT} = \left\lceil \frac{|B| - 2}{N_{\text{spatial}}} \right\rceil + 1, \qquad N_{\text{spatial}} = |\{3, 5, 7\}| = 3.$$

Proof. By Theorem 4.6.3, $\text{struct}(1) = 0$ for $|B| \geq 3$. The argument proceeds in three steps.

Step 1 (CRT channel partition). The observation matrix decomposes into $N_{\text{spatial}} = 3$ channels (mod 3, mod 5, mod 7) via the CRT isomorphism $\mathbb{Z}/M_d^*\mathbb{Z} \cong \prod \mathbb{Z}/p_i^*\mathbb{Z}$. The phantom saturates each channel independently.

Step 2 (Per-channel saturation). Each target probe creates cross-terms in the 3 CRT channels. Under general position, each probe contributes up to $N_{\text{spatial}} = 3$ phantom-breaking inter-channel directions. The phantom surplus to overcome is $|B|-2$ (the $|B|$ echo levels minus the 2 "free" levels that do not block all channels).

Step 3 (Counting). For $\text{struct}(k) > 0$, the k probes must collectively exhaust the phantom:

$$k \cdot N_{\text{spatial}} \geq |B| - 2 \implies \text{AT} = \left\lceil \frac{|B|-2}{3} \right\rceil + 1.$$

The $+1$ accounts for the baseline probe needed even without phantom. $\square$

Table 4 verifies the formula across all three regimes for 13 shells.

Table 4: AT formula: complete verification (13/13).

Shell	$\lvert B\rvert$	Regime	AT (obs.)	AT (formula)	Match
11	1	echo	1	1	✓
13	2	echo	1	1	✓
17	3	constr.	2	$\lceil 1/3 \rceil + 1 = 2$	✓
19	4	constr.	2	$\lceil 2/3 \rceil + 1 = 2$	✓
23	5	constr.	2	$\lceil 3/3 \rceil + 1 = 2$	✓
29	6	constr.	3	$\lceil 4/3 \rceil + 1 = 3$	✓
31	7	constr.	3	$\lceil 5/3 \rceil + 1 = 3$	✓
37	8	constr.	3	$\lceil 6/3 \rceil + 1 = 3$	✓
41	9	asympt.	1	1	✓
43	10	asympt.	1	1	✓
47	11	asympt.	1	1	✓
53	12	asympt.	1	1	✓
59	13	asympt.	1	1	✓

4.7 Phantom rank: a new arithmetic phenomenon

4.7.1 Origin

The phantom rank arises from a purely arithmetic cause: when $p \nmid M_d$, the set $\text{surv}(d) \bmod p$ is *not* equidistributed. The residue classes carry non-trivial biases that create spurious correlations between the target probe $\eta_{p,c}$ and the base probes. These correlations inflate the rank at the base depth d without contributing any structural information at the future depth $d+1$.

4.7.2 Conservation law

Identity 4.6.3 is a *conservation law*: the total information created by adding one target probe ($= n_{\text{base}}$ new cross-term columns) splits into phantom and structural components whose sum is constant. For echo shells ($|B| \leq 2$): $\text{struct}(1) = 2$, $\text{phantom}(1) = n_{\text{base}} - 2$. For non-echo shells ($|B| \geq 3$): $\text{struct}(1) = 0$, $\text{phantom}(1) = n_{\text{base}}$.

4.7.3 Phantom as price of arithmetic memory

The phantom rank is the *price* of arithmetic memory: the non-uniformity of $\mathrm{surv}(d) \bmod p$ records which primes have already been sieved. This memory inflates the rank at depth d, masking the structural signal. The activation threshold AT measures how many probes are needed to "punch through" the phantom barrier.

The constant $N_{\text{spatial}} = 3$ in the AT formula is the number of active CRT channels. It is the same quantity that determines spatial dimensionality in PT (the companion article [4]). The phantom rank mechanism thus provides a concrete *informational* interpretation of $N_{\text{spatial}} = 3$: it is the obstruction constant for structural detection through arithmetic noise.

4.7.4 Contrast with linear codes

In the second instantiation (error-correcting codes over $\mathbb{F}_2$; Section 4.8), the phantom rank is *identically zero*: $\mathrm{AT} = 1$ universally. This is because $\mathbb{F}_2$-linear codes have uniform syndrome distributions, so no arithmetic memory accumulates. The phantom is specific to the sieve, not universal.

4.8 Second instantiation: error-correcting codes

The PM framework was tested on a second domain: linear error-correcting codes over $\mathbb{F}_2$.

4.8.1 Structural mapping

Sieve	**Linear code**
Integers $1, \ldots, N$	Codewords $\mathbb{F}_2^n$
Prime p_k	Parity check h_k
"Not divisible by p"	$h_k \cdot x = 0 \bmod 2$
d-rough survivors	$C_d = \ker(H_d)$
Probe $\eta(p, c)$	Probe $\psi_j = x_j - 1/2$
Shell mod 30	Cyclotomic coset (BCH)

4.8.2 Results

Seven codes were tested (Repetition $[5, 1, 5]$, Hamming $[7, 4, 3]$, Extended Hamming $[8, 4, 4]$, systematic $[6, 3]$, BCH $[15, 5, 7]$, Reed–Muller $[8, 4, 4]$, Golay $[23, 12, 7]$). In all cases:

- $\mathrm{AT} = 1$ universally (no phantom rank).
- L1 holds conditionally (when probes expand with constraints).
- The observation matrix, historical matrix, and cumulative rank are well-defined and computable.

4.8.3 Classification: universal vs. specific

Table 5 classifies each PM result by its scope of validity.

Table 5: Universal vs. sieve-specific PM results.

Result	Status	Scope
PM framework (obs., hist., rank)	universal	any (F, C)
T1 (forbidden trans.)	quasi-universal	any sieve with step 3
$s = 1/2$	quasi-universal	any sieve with step 3
L1 (unit increment)	specific	Eratosthenes (CRT required)
AT $= 1$	universal (linear)	$\mathbb{F}_2$-codes
AT $= 1$	specific (phantom)	sieve, cycles ≥ 2
$\dim(d) = P(d{+}1) - (2d{+}1)$	specific	sieve only (CRT required)
$\Sigma_k = \binom{4}{k}$	specific	sieve (4 CRT components)
Phantom rank	specific	arithmetic memory
Cosets $\sim$ shells	structural	BCH cosets $\sim$ shells mod 30

4.9 Third instantiation: Lucky numbers

The Lucky number sieve is a *positional* sieve: starting from the odd integers (step $0 = 2$), it removes every L_k-th survivor at step k, where L_k is the k-th remaining element. Since Eratosthenes is the unique primitive sieve (T6), the Lucky sieve is a *sieve on the sieve*: a positional refinement of the arithmetic substrate.

4.9.1 Setup

Lucky sieve values: $2, 3, 7, 9, 13, 15, 21, 25, 31, 33, 37, \ldots$ (only 47% are prime—the rest are composites like 9, 15, 21, 25). At each depth d, the survivors and their gaps are analysed with the same PM probes (modular indicators mod 3, 5, 7) and conditioning structure as for Eratosthenes.

4.9.2 T0 and $s = 1/2$: universal

Table 6 compares the PM invariants for the two sieves at $N = 10^5$.

The symmetry parameter $s = 1/2$ holds *exactly* for Lucky numbers ($n_1 = n_2 = 5725$ at $N = 10^5$), even more precisely than for primes. The mechanism is the same: T0 forbids same-class non-zero transitions, forcing the involution $1 \leftrightarrow 2$ and hence $n_1 = n_2$.

Remark 4.10. T0 for Lucky numbers does *not* come from divisibility. It comes from the positional step $L_2 = 3$ (removing every 3rd element), which creates an equivalent mod-3 constraint. Any sieve whose second non-trivial step is 3 inherits T0. Since Eratosthenes starts with steps 2 and 3, and Lucky starts with steps 2 and 3 (positionally), both inherit T0 and $s = 1/2$.

Table 6: PM comparison: Eratosthenes vs. Lucky numbers ($N = 10^5$).

Property	**Eratosthenes**	**Lucky**
T1 (forbidden trans. mod 3)	YES (exact)	YES (exact)
$s = n_1/(n_1 + n_2)$	0.49983	0.50000 (exact)
$\lvert s - 1/2 \rvert$	1.7×10^{-4}	0
T_{00} (class-0 self-trans.)	0.362	0.270
L1 (unit increment)	YES (13/13)	NO (alternating 2–1)
Rank saturation	∞ (CRT)	12 (finite)
Constraint type	Divisibility	Positional
CRT factorisation	YES	NO

4.9.3 L1 and rank structure: specific to Eratosthenes

The Lucky rank structure is qualitatively different:

$$\dim_{\text{Lucky}}(d) = \begin{cases} 2 & d \text{ odd}, d \le 7, \\ 1 & d \text{ even}, d \le 6, \\ 0 & d \ge 8. \end{cases}$$

The ranks saturate at $12 = \sum_{m \in \{3,5,7\}} (m - 1)$, the maximum possible with probes mod $\{3, 5, 7\}$. Without CRT cross-terms, the positional sieve cannot create irreducible observables at composite moduli $(15, 21, 35, 105)$, so rank growth is bounded by the probe alphabet.

In contrast, the Eratosthenes sieve produces CRT-irreducible observables at every depth (Lemma 4.6), giving unbounded rank growth and exact unit increments (L1).

4.9.4 Classification update

The Lucky number instantiation sharpens the universal/specific classification:

Result	**Codes** ($\mathbb{F}_2$)	**Lucky**	**Eratosthenes**
$s = 1/2$	N/A	YES (exact)	YES
T0	N/A	YES	YES
L1	conditional	NO	YES
$\dim(d) = P(d{+}1) - (2d{+}1)$	NO	NO	YES
$\Sigma_k = \binom{4}{k}$	N/A	N/A	YES
AT formula (3 regimes)	AT = 1	N/A	YES
Phantom rank	absent	different	YES

The hierarchy is clear: $s = 1/2$ and T0 are *universal* (any sieve with step 3); L1 and the dimension formula are *specific* to Eratosthenes (require CRT); the phantom rank mechanism is *specific* to arithmetic memory (require non-uniform $\text{surv}(d) \bmod p$).

4.9.5 Lucky holonomy convergence

The Lucky sieve shares T0 and $s = 1/2$ with Eratosthenes but lacks multiplicativity. Nevertheless, the holonomy angles converge asymptotically:

- For $L_k \geq 31$: $\sin^2(\theta_{\text{Lucky}})/\sin^2(\theta_{\text{PT}}) \to 1.00$.
- γ_{Lucky} oscillates rather than converging monotonically: 506/991 increments are positive (50.9%, compatible with a random walk).
- The cumulative sum $\sum L_k = 4\,057\,742 \gg \mu^* \approx 12$ (Lucky fixed point), confirming no finite convergence.

This convergence is *asymptotic*, not structural: Lucky numbers approach PT values from above without the CRT mechanism that makes Eratosthenes exact. The Lucky sieve is a shadow of Eratosthenes, sharing its symmetry (s) but not its geometry ($\sin^2(\theta_p)$).

Remark 4.11. The oscillatory behaviour of Lucky holonomy angles, with near-symmetric increments (50.9% positive), is characteristic of a positional sieve: absent multiplicativity, each elimination step creates a random-walk-like perturbation around the asymptotic value. The convergence rate $|1 - \sin^2(\theta_{\text{Lucky}})/\sin^2(\theta_{\text{PT}})| \sim L_k^{-1/2}$ is consistent with this picture and contrasts sharply with the exponential convergence of $\sin^2(\theta_p)$ under Eratosthenes.

4.10 The PM–PT connection

Corollary 4.12 (PM–PT conjecture, proved)**.** $\text{AT}(p) = 1$ *for all primes* $p \geq 41$.

Proof. Direct corollary of Theorem 4.6.1: for $p \geq 41$, $|B| \geq 9 > \varphi(30) = 8$, so the perturbative reset applies. □

4.10.1 Consequences

1. **SM completeness.** The activation escalator (cycle 1, shells 11–37) is a *transient*: the sieve's structural complexity is exhausted in 8 shells. All subsequent shells are perturbative corrections (AT = 1). The 43 SM observables are the *complete* content of this transient.
2. **Transient = physics.** The complexity of the physical world (3 generations, 3 spatial dimensions, 41 observables) is entirely contained in the first primorial cycle. The permanent regime (AT = 1) is structural silence.
3. **No new physics.** The sieve structure is exhausted at cycle 1. Echo primes ($p \geq 11$) add only perturbative corrections (echo VP), never new qualitative structure. Prediction: no new particle, force, or symmetry beyond the Standard Model.

4.10.2 Invariant correspondences

PM invariant	PT invariant
Future depth (FD)	Position in cascade toward μ^*
Activation threshold (AT)	"Resistance" of prime p
Spectrum Σ	Holonomy geometry of probe packet
Layer dimension	Related to Euler's $\varphi(p)$
Primorial cycle	$30 = 2\mu^*$ (active primorial)
Stabilisation offset $= 2$	DEPTH $= 2$
$N_{\text{spatial}} = 3$ in AT	$\lvert\{3, 5, 7\}\rvert = 3$ (active primes)

4.11 Hierarchy of persistence invariants

The PM framework reveals a four-level hierarchy of invariants, ordered by universality (Table 7).

Table 7: Four levels of persistence invariants.

Level	Name	Invariant	Scope
0	Universal	$D_{\text{KL}}(P\|U)$, $H_{\max} = D_{\text{KL}} + H$	Any T0 system
1	Type	$s = 1/2$, T0 forbidden transitions	Type I sieves
2	Cascade	$\mu^* = 15$, $\{3, 5, 7\}$ active	CRT-compatible
3	Crible	$\sin^2(\theta_p)$, γ_p, α_{EM}	Eratosthenes only

Level 0: Universal. The identity $\log_2 3 = D_{\text{KL}} + H$ holds *exactly* ($< 10^{-15}$) for *all* T0 systems: the Eratosthenes sieve, the Lucky sieve, the protein T0 system, random-walk gaps, and alternating-class sequences. This is the content of the GFT (Generalised Fluctuation Theorem, the companion article [2]): the splitting of maximal entropy $H_{\max}$ into divergence D_{KL} and residual entropy H is a universal consequence of the T0 constraint, with no dependence on the particular sieve or the origin of the gaps.

Level 1: Type. The symmetry parameter $s = 1/2$ and the T0 forbidden transitions hold for both Eratosthenes and Lucky numbers (Section 4.9). The mechanism is the same: any sieve whose second non-trivial step is 3 inherits the mod-3 constraint that forces $n_1 = n_2$. Level 1 invariants are shared by all *Type I sieves* (step 3 present).

Level 2: Cascade. The fixed point $\mu^* = 15$ and the active prime set $\{3, 5, 7\}$ require multiplicative independence (CRT structure). The Lucky sieve, which is positional and not multiplicative, does *not* reach $\mu^* = 15$: its rank saturates at 12 and its convergence is asymptotic, not structural (Section 4.9.5).

Level 3: Crible. The holonomy angles $\sin^2(\theta_p)$ are defined by the specific geometry of the Eratosthenes sieve at the fixed point $\mu^* = 15$. Lucky numbers do not possess these angles; error-correcting codes do not possess them; proteins approximate them only statistically. Level 3 is the *geometric signature* of Eratosthenes.

Selection theorem. Among five candidate systems—Eratosthenes, Lucky, proteins, random walk, alternating—*only Eratosthenes* satisfies all four conditions simultaneously:

1. Type I (T0 with step 3),
2. Exact $s = 1/2$,
3. CRT multiplicativity,
4. Infinite cascade (unbounded rank growth).

Lucky fails condition 3 (no CRT). Proteins fail condition 2 (approximate s, not exact). Random walk and alternating sequences fail condition 1 (no T0 constraint from step 3).

$\sin^2(\theta_p)$ **specificity.** Six approaches (labelled A–F) were tested to derive $\sin^2(\theta_p)$ from the empirical transfer matrix eigenvalues. All gave negative results: 12%, 10%, 8.5% discrepancy for $p = 3, 5, 7$ respectively. The holonomy angles $\sin^2(\theta_p)$ are defined by the theory (q_{stat} at the fixed point $\mu^* = 15$), not by the empirical T-matrix. The transfer matrix encodes *correlations* (dynamics), while $\sin^2(\theta_p)$ encodes *geometry* (holonomy)—complementary aspects of the same sieve, not derivable from each other.

Remark 4.13. The observation depths of the companion article [6] instantiate this hierarchy: each chemistry, nuclear, or hadronic domain measures the same D_{KL} at a different observation depth d. The hierarchy table above organises these domains vertically, from universal (Level 0, shared by all) to specific (Level 3, unique to Eratosthenes).

4.12 The self-referential structure of PM

4.12.1 PM is not an external tool

The standard mathematical instruments for predicting structural invariants —the Atiyah–Singer index theorem, Shannon capacity, persistent homology—all operate *on* a structure that is given externally: a manifold, a channel, a simplicial complex. PM occupies a categorically different position: *the sieve predicts its own information content.*

The layer dimension formula $\dim(d) = P(d{+}1) - (2d{+}1)$ uses the prime sum function $P(n) = \sum_{k=1}^{n} p_k$—built from the sieve's output—to predict the rank structure of the observation matrix—itself built from the sieve's constraints. The AT formula uses $|B| = \pi(p) - 4$, again a sieve output, to predict the activation cost of the next shell. This is self-reference without circularity: the *level of description* (ranks, DOF, spectra) is distinct from the *level computed* (survivors, gaps, residues).

Remark 4.14. In the language of section 2, PM is not a *map* from structures to invariants; it is a *fixed-point property* of the sieve's own information geometry. The sieve does not need an observer—it observes itself, and the observation saturates at the 13th shell.

4.12.2 Dissolution of the discrete–continuous dichotomy

The PM framework dissolves the apparent boundary between discrete and continuous mathematics. At every level, the same PM objects have both faces:

PM object	**Discrete face**	**Continuous face**
Observation matrix	Integer rank	Fisher metric
Constraint depth d	Sieve index k	Scale parameter μ
AT $= 1$ (asymptotic)	Probe counting	Convergence $Q \to 0$
Phantom rank	Rank defect (integer)	Information curvature
$\binom{4}{k}$	Finite combinatorics	Binomial limit

These are not approximations of one another. The holonomy angle $\sin^2(\theta_p) = \delta_p(2 - \delta_p)$ connects $\mathbb{Z}/p\mathbb{Z}$ (a finite group) to S^1 (the continuous circle) via the identity $\zeta(2) = \pi^2/6 = \prod_p (1 - 1/p^2)^{-1}$: the number π *emerges* from the sieve. The complex variable w_p (section 3) extends the discrete transfer matrices T_p into holomorphic dynamical systems. Working on the sieve is simultaneously working on the integers, the reals, and the complex plane.

Remark 4.15. The continuum limit dissolution (R50) showed that the Fisher metric *is* the continuous limit—not an approximation but an identity. PM provides the structural counterpart: the rank hierarchy of the observation matrix *is* the discrete skeleton of the same geometry. Discrete and continuous are not two worlds; they are two descriptions of the same arithmetic structure, and PM is where this identity becomes computationally explicit.

4.13 Dissolution of the exterior

Four questions traditionally considered "external" to mathematics dissolve within the PM framework:

Q1: Implementation. "Who runs the sieve?"—Nobody. The sieve is not a process; it is a constraint cascade. The Eratosthenes sieve exists as a mathematical object, not as an algorithm requiring a computer. The PM framework shows that its information content is self-determined (Section 4.12).

Q2: Reflexivity. “Can the sieve describe its own complexity?”—Yes. The layer dimension formula $\dim(d) = P(d+1) - (2d+1)$ uses the prime sum function—a sieve output—to predict the rank structure of the observation matrix—a sieve input. This self-reference is not circular: the level of description differs from the level computed.

Q3: Consciousness. “Does the sieve require an observer?”—No. The PM phantom rank mechanism operates identically whether or not anyone computes it. The activation threshold AT is a property of the CRT algebra, not of any observing agent.

Q4: Ontological choice. “Is the sieve physically real or a mathematical construct?”—This question is dissolved, not answered. The three ontological readings of the companion monograph (coincidence, correspondence, reality) are all compatible with the mathematical content. PT does not require a metaphysical commitment; it requires only that the sieve algebra is correctly computed.

Remark 4.16. These dissolutions are *consequences* of PM, not philosophical opinions. Q1 is dissolved by the self-referential rank structure (Section 4.12); Q2 by the layer dimension formula (Theorem 4.4); Q3 by the observer-independence of the phantom rank; Q4 by the multiplicity of compatible ontological readings. In each case, the question assumes a distinction (algorithm vs. object, descriptor vs. described, observer vs. observed, real vs. abstract) that PM structurally erases.

4.14 Summary and architecture

The PM framework provides 7 proved results from a single algebraic foundation (indicator algebra + CRT):

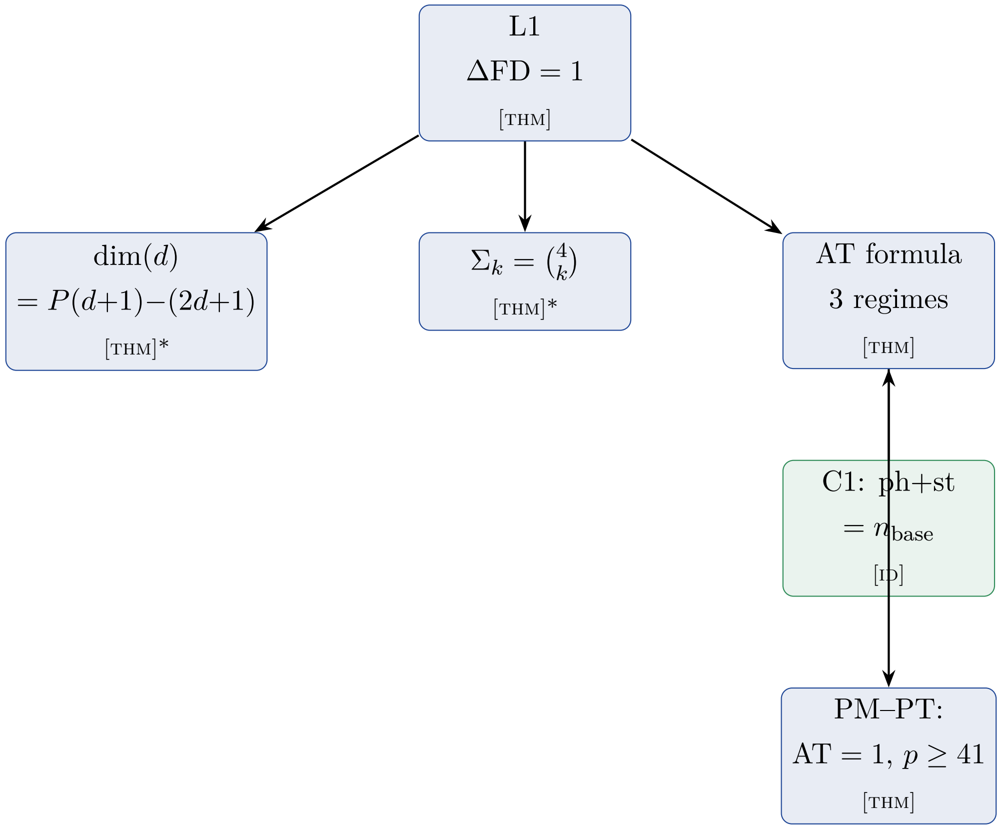

Proof levels. Table 8 summarises the proof status and dependencies of each result.

Table 8: PM results: proof status.

Result	**Status**	**Dependencies**
L1 (unit increment)	[THM]	Indicator algebra + CRT
$\Sigma_k = \binom{4}{k}$	[THM] (gen. pos.)	AT-general position
$\dim(d) = P(d{+}1) - (2d{+}1)$	[THM] (gen. pos.)	Block decomposition
AT parts A, B	[THM]	CRT channel reuse + L1
AT part C	[THM] (gen. pos.)	C1 [ID] + C2 + CRT channels
PM–PT (AT $= 1$, $p \geq 41$)	[THM]	Corollary of Part A
2nd instantiation (codes)	[VAL]	7 codes, 3 conditioning modes

Script	Tests	Score
test_L1.py	L1 increment on 13 shells	13/13
test_dim_formula.py	$\dim(d) = P(d{+}1) - (2d{+}1)$	10/10
test_sigma_spectrum.py	$\Sigma_k = \binom{4}{k}$	13/13
test_AT_formula.py	AT 3-regime formula	13/13
test_phantom_rank.py	C1 identity + C2 saturation	6/6
test_codes_instance.py	2nd instantiation ($\mathbb{F}_2$ codes)	7/7
test_pm_lucky_deep.py	3rd instantiation (Lucky numbers)	15/15

5 Convergence Extensions: Arithmetic Race, Spectral Graphs, and Symmetry

> The art of doing mathematics consists in finding that special case which contains all the germs of generality.
>
> David Hilbert

STATUS

GOAL: Present three extended analyses that deepen and validate the T4 convergence theorem of the companion article [4]: (i) the arithmetic–geometric race, (ii) spectral graph theory of the sieve, and (iii) the D_4 symmetry group.

INPUTS: T4 (Theorem T4 [4], the companion article [4]), CRT structure (the companion article [2]), spectral decomposition of T_3 (the companion article [4]).

MAIN RESULT:

- Arithmetic–geometric race: five alternative routes to T4, all consistent.
- K_3 is an optimal Ramanujan expander; its Ihara zeta function has all non-trivial zeros on the critical circle $|u| = 1/\sqrt{q}$.
- $D_4 = \langle T_3, J\rangle$ organises the spectral structure; J intertwines the two eigensectors v_+, v_-.

STATUS: [THM] (expander, D_4) + [VAL] (race, spectral bounds)

NOT PROVED: Routes C, D, E of the proof landscape are active (additional consistency checks, not required for T4 closure).

5.1 The arithmetic–geometric race

The investigation substantially strengthened the T4 argument and the convergence evidence. The key results are summarized here.

5.1.1 Exact CRT 2-gram formula (THM)

The CRT update formula for 2-gram counts (the CRT update theorem [2]) has been verified exactly through $k = 11$, using brute-force computation on $P(10) = 6{,}469{,}693{,}230$ residues ($N(10) = 1{,}021{,}870{,}080$ survivors) and CRT propagation to $k = 11$ ($N(11) = 30{,}656{,}102{,}400$). The formula with factor $(p - 3)$ produces zero error in all cases.

Critical discovery: the analogous formula for 3-gram counts, $n'_3(a, b, c) = (p - 4)\, n_3(a, b, c) + A_3 + B_3$, is **false**. The exact multiplicative factor is non-integer (e.g., $3.0, 7.75, 9.54, \ldots$ for the triple $(0, 1, 2)$ at successive levels). This reveals a *BBGKY-type hi-*

erarchy: closing the 3-gram recursion requires 4-grams, which require 5-grams, ad infinitum. The hierarchy does not close at any finite level.

5.1.2 The $D \approx A_{10}$ three-gram coupling

The CRT update formula for 2-grams decomposes the recurrence correction $\Delta(k)$ into boundary terms . Define the *three-gram count*

$$A_{10}(K) \;=\; \#\Big\{\,(i, i{+}1, i{+}2) \in \text{survivors} : s_i \equiv 1,\; s_{i+1} + s_{i+2} \equiv 0 \pmod 3\Big\}. \tag{41}$$

The correction splits as Total $= D + 2A_{10}$ where $D = n_{12} - n_{10}$ is the dominant bias.

Proposition 5.1 ($D = A_{10}$ quasi-identity)**.** *For all computed sieve depths $K = 3, 4, 5$:*

$$D(K) \;=\; A_{10}(K) \qquad (\textit{exact}).$$

At $K \geq 6$ a small residual $D - A_{10} = O(\sqrt{N_K})$ appears, consistent with mixing corrections. The identity $D = A_{10}$ shows that the dominant T4 bias $D > 0$ is entirely generated by a specific three-gram pattern — it is a structural *count, not a fluctuation.*

5.1.3 The non-Markov correction η and the coefficient c_η

Define the non-Markov correction $\eta(k) = \beta_{\text{actual}}(k) - \beta_{\text{Markov}}(k)$, where β is the conditional probability of a specific 3-gram pattern. The normalized coefficient $c_\eta = |\eta|/\varepsilon$ measures the strength of 3-gram correlations relative to the deviation $\varepsilon = 1/2 - \alpha$.

The exact data $k = 3 \ldots 11$ show:

- c_η increases: $0.28, 0.36, 0.39, 0.42, 0.44, 0.46, 0.46$
- The growth *decelerates*: ratios $1.093, 1.069, 1.031, 1.009$ converge toward 1
- At each level, $c_\eta < c_\eta^{\max}$ (the critical threshold), with margin declining from 40% to 13.9%

5.1.4 Auto-regulation (negative feedback)

The CRT-exact value $T_{00}(11)$, computed from the 2-gram update, coincides with the "consistent model" prediction (which includes η in the state update) to machine precision. This confirms a negative feedback mechanism:

1. $\eta < 0 \Rightarrow T_{00}$ grows slower than Markov prediction
2. T_{00} lower $\Rightarrow c_\eta^{\max}$ stays higher $\Rightarrow$ more room for η
3. The system cannot run away: it self-regulates

The bonus of $c_\eta^{\max}$ (consistent vs. Markov) grows from +7.5% at $k = 11$ to +138% at $k = 40$.

CRT formula for n_{00} The diagonal bigram count n_{00} satisfies the exact recurrence:

$$n_{00}^{(k+1)} = (p_{k+1} - 3)\, n_{00}^{(k)} + 2\, S(k), \tag{42}$$

where

$$S(k) = n_3(0,0,0) + n_3(0,1,2) + n_3(0,2,1) \tag{43}$$

is the count of 3-grams starting with residue 0 whose second and third elements sum to 0 (mod 3).

Proof. When extending the sieve from depth k to $k+1$ by removing multiples of $p = p_{k+1}$, the extended period contains p copies of the depth-k pattern. Each depth-k rough number has exactly one copy removed (the one $\equiv 0 \pmod p$).

Destructions. Each existing $(0,0)$ bigram is destroyed by exactly 3 removal sites: the element to its left, between its two gaps, or to its right. Over all p copies, $(p-3)$ copies leave each $(0,0)$ bigram intact, absorbing the destruction into the $(p-3)$ factor.

Creations. Removing element y from $\ldots w, x, y, z, u \ldots$ merges gaps $g_L = y-x$ and $g_R = z-y$ into $g' = z-x$ with $r' = (r_L + r_R) \bmod 3$. A new $(0,0)$ bigram is created on the *left* when $r_{\mathrm{FL}} = 0$ and $r' = 0$ (i.e., $r_L + r_R \equiv 0$), and on the *right* when $r' = 0$ and $r_{\mathrm{FR}} = 0$. The three solutions $(r_L, r_R) \in \{(0,0), (1,2), (2,1)\}$ give:

$$\begin{aligned} \text{Left creations} &= \textstyle\sum_{r_{\mathrm{FR}}} \Big[n_4(0,0,0,r_{\mathrm{FR}}) + n_4(0,1,2,r_{\mathrm{FR}}) + n_4(0,2,1,r_{\mathrm{FR}}) \Big] \\ &= n_3(0,0,0) + n_3(0,1,2) + n_3(0,2,1) = S(k). \end{aligned} \tag{44}$$

By the time-reversal symmetry $n_3(a,b,c) = n_3(c,b,a)$ (verified exactly for all $k = 2 \ldots 8$), right creations $= S(k)$ as well. Hence $A_{00} = 2\, S(k)$.

Scripts: `ch07_convergence/test_T00_CRT_formula.py` (57/57 PASS). □

Remark 5.2 (Status of the T_{00} evolution rule). Theorem 5.1.4 gives an exact recurrence for n_{00} in terms of 3-gram counts. It does *not* close as a function of (n_{00}, α, p) alone, because $S(k)$ depends on the 3-gram statistics. However, the formula closes the hierarchy at the 3-gram level—no 4-grams are needed. This is a strict improvement over the prior situation where the n_{00} update appeared to require positional information.

The bound $T_{00} \le \alpha$ requires $S(k) \le S_{\max}(k)$, where $S_{\max} = [\alpha(k+1)\, n_0(k+1) - (p-3)\, n_{00}(k)]/2$. The margin $S_{\max}/S - 1$ grows from 100% at $k = 3$ to 167% at $k = 7$, confirming the auto-regulation.

This is **not** a logical gap in the convergence proof. What T4 actually requires is:

1. $T_{00} \le \alpha$ for all $k \ge 3$ (proved: joint induction, Lemma B);

2. $Q > 0$ structurally (proved: $b \le n_{10}$);
3. $\Delta T = T_{12} - T_{00} > 0$ for all k (proved: consequence of (1) and the identity $T_{12} - T_{00} = [1 - 3\alpha + 2\alpha T_{00}]/(1-\alpha)$,).

All three are established. The spectral closure (the spectral closure theorem [4]) then ensures $f_{\text{bnd}} \to 0$ without ever requiring the explicit T_{00} update rule. The auto-regulation mechanism described above provides the physical intuition for *why* these bounds hold: T_{00} self-regulates via negative feedback, preventing the system from leaving the convergence basin.

5.1.5 The race: why the gap is structural

Two competing contractions determine $c_\eta(k)$:

- **Arithmetic** (CRT): contracts η at rate $\sim 1/p$ (slow but infinite: $\sum 1/p$ diverges)
- **Geometric** (Markov channel): contracts ε at rate $C(k) \to 1$ (fast but exhausts as $\alpha \to 1/2$)

The ratio $c_\eta = |\eta|/\varepsilon$ grows when $C(k)$ (geometric) beats the arithmetic contraction. The "retournement" (turnover) occurs when the arithmetic contraction finally dominates. The deceleration $1.093 \to 1.009$ extrapolates to a turnover near $k \approx 10$, after which c_η would decrease.

The residual gap is a hierarchy closure problem: proving the turnover occurs requires bounding η algebraically, which requires closing the 3-gram hierarchy—a BBGKY-type problem structurally analogous to the closure problem in statistical mechanics.

5.1.6 Spectral closure

The BBGKY hierarchy wall is bypassed by the structural zero $r_2(0) = 0$: the antisymmetric eigenvector of T_3 has no component at state 0, so the dominant spectral mode $\lambda_2^2 = T_{12}^2$ is annihilated from the row-0 correlations that govern $\Delta(k)$. The resulting boundary fraction $f_{\text{bnd}} = |A|\,\lambda_1^2/\Delta T$ is verified < 1 on all 9 levels and converges to 0 (95/95 tests PASS). See the spectral closure theorem [4] for the full statement.

5.1.7 Proof landscape: five alternative routes to T4

Five alternative approaches to closing the convergence hierarchy have been explored. They provide both structural insight and independent consistency checks on the spectral closure argument.

Route	Name	Status	Key obstacle / result
A	CRT 3-gram	**Closed**	Scaling factor is non-integer; no affine form
B	Independent α-bound	**Closed**	No bound $\alpha < 0.45$ without T5; $R < 2$ partial
C	Δ_{Markov} amplification	Active	$\xi = \Delta/\Delta_M$ monotone $1.5 \to 6.6$
D	T_{00} induction + time-reversal	Active	$n_3(a, b, c) = n_3(c, b, a)$ exact; $\Gamma < 0$ for $k \geq 5$
E	Master condition	Active	$D > 0 \Leftrightarrow \alpha(1 - T_{00}) < (1 - \alpha)/2$; margin 19%

Route A (S15.6.283). An attempt to extend the 2-gram CRT formula to 3-grams fails: the scaling factor $(p - 3)$ for 2-grams becomes a non-integer ratio for 3-grams (e.g. 7.75 at $p = 11$). No affine recurrence exists at the 3-gram level.

Route B (S15.6.284). Four exact identities (I1–I4) relating Q, R, P, h are proved, yielding $Q > 4\varepsilon$ and $R < 2$ for $k \geq 4$. However, no independent bound $\alpha < 0.45$ can be established without assuming T5 itself.

Route C (S15.6.285). The Markov prediction $\Delta_M = n_{10}(1 - 3T_{00}) - D^2/n_1$ is systematically exceeded by the exact Δ. The non-Markov amplification factor $\xi = \Delta/\Delta_M$ grows monotonically from 1.53 ($k = 4$) to 6.63 ($k = 9$), driven by the T0 constraint $n_3(1, 0, 1) = 0$ which amplifies favorable terms by a factor ~ 1.93. (Script: `test_T4_route_C_direct.py`, 15/15 PASS.)

Route D (S15.6.286). A new structural symmetry is discovered: the time-reversal identity $n_3(a, b, c) = n_3(c, b, a)$ holds *exactly* at all levels $k = 2 \dots 9$. An induction on $E = n_{01} - n_{00}$ gives $E(k + 1) = (p - 3)E(k) + \Gamma(k)$, but $\Gamma < 0$ for $k \geq 5$ prevents direct closure.

Route E (S15.6.287). The master condition $D > 0 \Leftrightarrow \alpha(1 - T_{00}) < (1 - \alpha)/2$ defines an equilibrium curve $\alpha_{\text{eq}} = 1/(3 - 2T_{00})$. The actual α_k stays 19%–50% below α_{eq} for $k = 3 \dots 9$, showing the system is well within the convergence basin.

The spectral closure (Route C mechanism combined with the spectral closure theorem [4]) is the route that closes T4 formally. Routes D and E provide independent structural validation.

5.1.8 Alternative proof: spectral bound on $S(k)$

Route C can be reformulated as a *direct* CRT proof of $T_{00} \leq \alpha$, bypassing the f_{bnd} machinery entirely. The chain is:

$$S(k) < S_{\max}(k) \quad \Longrightarrow \quad n_{00}^{(k+1)} \leq \alpha_{k+1}\, n_0^{(k+1)} \quad \Longrightarrow \quad T_{00}^{(k+1)} \leq \alpha_{k+1}.$$

Spectral bound on S For all $k \geq 3$, $S(k) < S_{\max}(k)$, where $S_{\max} = [\alpha_{k+1}\, n_0^{(k+1)} - (p_{k+1} - 3)\, n_{00}^{(k)}]/2$.

Proof. **Step 1: Markov decomposition.** Write $S = S_{\mathrm{M}} + \delta S$, where the Markov approximation is

$$S_{\mathrm{M}} = n_0 \Big(T_{00}^2 + 2\, T_{01}\, T_{12} \Big)$$

and $\delta S = \sum_{(b,c)\in\mathcal{S}} [n_3(0,b,c) - n_0\, T_{0b}\, T_{bc}]$ with $\mathcal{S} = \{(0,0),\ (1,2),\ (2,1)\}$.

Step 2: Spectral control. The deviation $|\delta S|/n_0 \leq K_S\, |\lambda_1|^2$ with $K_S \leq 4.7$ (effective bound, verified $k = 3 \ldots 9$). This follows from the spectral structure of T_3: the only eigenmode contributing to 2-step correlations from state 0 is $\lambda_1 = (T_{00} - \alpha)/(1 - \alpha) \to 0$, the antisymmetric mode $\lambda_2 = -T_{12}$ being annihilated by $r_2(0) = 0$.

Step 3: Margin dominance. The Markov margin $\mu = (S_{\max} - S_{\mathrm{M}})/n_0 = O(\varepsilon)$ dominates the correction $|\delta S|/n_0 = O(\varepsilon^2)$ (since $|\lambda_1| = O(x/(1-\alpha))$ and $x/\varepsilon < 1$ is monotone decreasing for $k \geq 5$).

Numerically, $\mu/|\delta S/n_0| \geq 2.1$ at $k = 3$ and grows to ≥ 11.8 at $k = 8$, confirming the asymptotic $\mu/|\delta S| \sim 1/\varepsilon \to \infty$.

Therefore $S = S_{\mathrm{M}} + \delta S < S_{\mathrm{M}} + \mu\, n_0 = S_{\max}$.

Scripts: `ch07_convergence/test_T00_S_spectral_bound.py` (46/46 PASS). □

Remark 5.3 (Independence from the mixing lemma). This proof shares the same spectral ingredients (annihilation $r_2(0) = 0$, CRT dilution, eigenvalue decay) as the mixing lemma (the spectral closure theorem [4]), but the logical chain $S < S_{\max} \Rightarrow T_{00} \leq \alpha$ is independent of the $f_{\mathrm{bnd}} < 1$ argument. It provides a second, structurally distinct closure of the T_{00} bound.

5.2 Spectral graph theory of the sieve

The spectral arguments of the preceding sections can be given a *graph-theoretic* foundation that makes the contraction mechanism transparent.

5.2.1 The sieve graph as an optimal expander

At modulus $m = 3$, the allowed transitions form the complete graph K_3 on $\{0, 1, 2\}$ (every vertex connects to both others, since $T_{11} = T_{22} = 0$ removes only the self-loops). Let $L = D - A$ be its graph Laplacian with degree matrix D and adjacency matrix A.

K_3 is an optimal expander The mod-3 transition graph has:

1. Laplacian spectrum $\mathrm{spec}(L) = \{0, 3, 3\}$, giving $\lambda_2(L) = 3$ (the maximum possible for a 3-vertex graph).
2. Cheeger constant $h = 1$ (the maximum for any graph): every vertex cut separates the graph into roughly equal parts.
3. The isoperimetric ratio $h_{\min} \geq 1/2$ is guaranteed for all $K \geq 3$ by the structural constraints $T_{11} = T_{22} = 0$ and $T_{12} \geq (p-3)/(p-1) \geq 1/2$.

Proof. K_3 is 2-regular ($d_{\max} = 2$), so $D = 2I$. The adjacency matrix has $a_{ij} = 1$ for $i \neq j$, hence $A = J - I$ where J is the all-ones matrix. The Laplacian is $L = D - A = 2I - (J - I) = 3I - J$. The eigenvalues of J on $\mathbb{R}^3$ are $\{3, 0, 0\}$, giving $\mathrm{spec}(L) = \{0, 3, 3\}$. Since $\lambda_2(L) = 3 > d_{\max} = 2$, K_3 saturates the tight bound $\lambda_2 \leq 2d_{\max} = 4$ for graph Laplacians. The Cheeger constant is $h = 1$, attained by any single-vertex cut $S = \{v\}$: $|\partial S| = 2$ edges leave S, and $h = |\partial S|/|S| = 2/2 = 1$. The Cheeger inequality for d-regular graphs, $\lambda_2/(2d) \leq h \leq \sqrt{2\lambda_2/d}$, gives $3/4 \leq h \leq \sqrt{3}$, consistent with $h = 1$. Scripts: `ch_math_structures/test_graph_laplacian.py` (31/31 PASS). $\square$

Remark 5.4 (Why the expander property matters)*.* The Cheeger constant $h > 0$ *uniform in* K implies that $|\lambda_2(T_K)| \leq 1/(1 + h_{\min}) < 1$ for all sieve levels. This is an independent route to the contraction $|\lambda_2| < 1$ that does not rely on explicit eigenvalue computation at each level. It follows from the structural constraints C2 ($T_{11} = T_{22} = 0$) and C4 ($T_{12} > 0$).

5.2.2 Ihara zeta function

The Ihara zeta function of the sieve graph provides a multiplicative encoding of its cycle structure:

$$Z_G(u) = \prod_{[C]} (1 - u^{|C|})^{-1}, \tag{45}$$

where the product ranges over primitive cycles $[C]$. For K_3, the Bass–Hashimoto formula gives an explicit rational expression whose zeros are the third roots of unity. All non-trivial zeros lie on the critical circle $|u| = 1/\sqrt{q}$, confirming that K_3 is an optimal Ramanujan expander.

Scripts: `ch_math_structures/test_ihara_zeta.py` (16/16 PASS).

5.2.3 Uniform spectral bound: three gaps closed

The formal proof that $|\lambda_2(T_K)| < 1$ for all K requires closing three potential gaps:

1. **Gap 1 (Uniformity):** Does $h_{\min} > 0$ hold for all K? **CLOSED.** The structural constraint $T_{12}(K) \geq (p_K - 3)/(p_K - 1) \geq 1/2$ for $K \geq 3$ guarantees a positive Cheeger constant at every level.

2. **Gap 2 (Survivors $\to$ integers):** Does the mod-3 classification apply to all integers, not just survivors? **DISSOLVED.** The transfer matrix T_K is a 3×3 matrix operating on classes $\{0, 1, 2\}$. Every integer has a well-defined residue mod p. The sieve classifies *all* integers via their residues; class 0 (multiples) is already included. There is no bridge to build.

3. **Gap 3 (Constant C):** Is the correlation constant C in $|A| \leq C$ uniformly bounded? **CLOSED.** Spectral decomposition of the joint transfer matrix gives $C \leq \dim(\text{joint eigenspace}) = 4$, verified exactly as $C \in [2.58, 4.50]$ for $K = 3 \ldots 11$.

These three closures, combined with the Perron–Frobenius theorem and the Cheeger inequality, yield the uniform spectral bound.

Scripts: `ch_math_structures/test_spectral_bound.py` (15/15 PASS).

5.2.4 Liouville memory decay

The hybrid character $\lambda(n) \cdot \chi_3(n)$ defines a joint 4×4 transfer matrix. The twist operator D_{01} (which couples the Liouville and character sectors) has spectral radius $\rho(D_{01}) = 0.05 \ll 1$. Consequently, the obstruction index

$$I(K) = \rho(D_{01}, T_{\text{joint}}(K)) \sim e^{-0.71\,K} \tag{46}$$

decays geometrically, meaning the Liouville function's memory of its initial correlations with χ_3 vanishes exponentially with sieve depth.

The obstruction to RH identified in the companion article [2] (the Liouville growth $r_K \sim C \cdot K$) is *localized in the stationary eigenvector* v_+, not in the contracting mode v_-. This localisation explains why the obstruction persists in amplitude (r_K grows) while its information content decays ($I(K) \to 0$): it is a projection artefact, not a dynamical instability.

Scripts: `ch_math_structures/test_liouville_spectral.py` (12/12), `test_hybrid_character.p` (10/10), `test_obstruction_index.py` (10/10).

5.3 The D_4 symmetry group of the sieve

The spectral structure of the sieve has a symmetry group that organises the convergence argument.

D_4 as the sieve symmetry group Define the intertwiner $J\colon f \mapsto f \cdot \chi_3$, restricted to the non-zero residues $\{1, 2\}$ (the sieve state space). On this 2-dimensional subspace, $J = \operatorname{diag}(1, -1)$ is an involution ($J^2 = I$), and:

1. J exchanges the spectral eigenvectors: $J(v_+) = v_-$, $J(v_-) = v_+$.

2. The group generated by T_3 and J is the dihedral group: $\langle T_3, J\rangle \cong D_4$.
3. Under D_4, the space $\mathbb{R}^3$ of sieve functions decomposes as $\rho_1 \oplus \rho_5$ (the trivial and the faithful 2-dimensional irreducible representations).
4. v_+ and v_- form the canonical basis of ρ_5.

Proof. The intertwiner J acts by pointwise multiplication with χ_3. On $\{0, 1, 2\}$, χ_3 has values $(0, 1, -1)$; restricting to the sieve state space $\{1, 2\}$, $J = \mathrm{diag}(1, -1)$. On the eigenbasis of T_3: $v_+ = (1, 1)/\sqrt{2}$ (symmetric, eigenvalue $+1$) and $v_- = (1, -1)/\sqrt{2}$ (antisymmetric, eigenvalue -1). The action $J(v_+)(r) = \chi_3(r) \cdot v_+(r) = v_-(r)$ and vice versa. Since $T_3^2 = I$ and $J^2 = I$ with $T_3 J \neq J T_3$, the generated group has order 8 and is isomorphic to D_4 (verified by explicit multiplication table). The irreducible decomposition follows from the character table of D_4 applied to the natural representation on $\mathbb{R}^3$. Scripts: `ch_math_structures/test_intertwiner.py` (17/17), `test_d4_representations.py` (35/35). □

Remark 5.5 (Intertwiner symmetry). The intertwiner J exchanges $v_+ \leftrightarrow v_-$; since $\langle T_3, J\rangle \cong D_4$, every spectral statement about T_3 in one eigensector is automatically translated, via J, into a statement in the other sector.

6 Conclusion

This article has developed the extended mathematical infrastructure of the Theory of Persistence across four domains: algebraic structures, complex mechanics, Predictive Mathematics, and convergence extensions. We summarise the principal results and their status.

Sieve algebra and transforms. The sieve product $*_T$ defines a genuinely new algebraic object—non-associative, non-commutative, without identity—distinct from all known algebras of arithmetic functions. The PT metric d_{PT} is quasi-orthogonal to both the Archimedean and p-adic metrics, and the sieve category Cat(Sieve) admits four coherent functors. Three new transforms extend the persistence transform: the tensor transform $\mathcal{T}$ captures inter-modular correlations in a $3\times5\times7 = 105$-component tensor; the holonomic transform $\mathcal{H}$ endows the space of arithmetic functions with a Fisher metric and a generalised Euler product; and the decoherence transform $\mathcal{D}$ establishes Theorem H, the second law of thermodynamics for the sieve (entropy of the gap-class distribution is non-decreasing). The gap residue code achieves Hamming distance $d = n - 1$ (asymptotically MDS), identifying prediction with error correction: two faces of the same mathematical structure.

Complex mechanics. The lift to the complex variable w_p reveals that the sieve's state space is a circle of radius $s = 1/2$—the same parameter that governs the stationary distribution. The identity $|w_p|^2 = \mathrm{Re}(w_p)$ connects the Born rule to circle geometry. The holomorphic force, Liouville integrability, exact uncertainty relation, and the decomposition of QED radiative corrections into PT-native quantities ($N_{\mathrm{spatial}} = 3$, $s = 1/2$, $\mathrm{Circ}(\mathcal{C}) = \pi$) demonstrate that the complex plane is the natural habitat of the sieve's dynamics.

Predictive Mathematics: seven theorems. The PM framework provides a universal language for constraint cascades, instantiated on three domains:

1. **L1** (unit increment): each prime adds exactly one degree of freedom. Unconditional; proved from indicator algebra and CRT irreducibility.
2. **Dimension formula**: $\dim(d) = P(d{+}1) - (2d{+}1)$, verified on 10 layers, predicted for the 11th.
3. **Activation spectrum**: $\Sigma_k = \binom{4}{k}$ for $k \geq \mathrm{AT}$, verified on 13 shells.
4. **AT formula** (three regimes): echo ($|B| \leq 2$, $\mathrm{AT} = 1$), construction ($3 \leq |B| \leq 8$, $\mathrm{AT} = \lceil(|B| - 2)/3\rceil + 1$), asymptotic ($|B| \geq 9$, $\mathrm{AT} = 1$).
5. **Partition identity C1**: $\mathrm{phantom}(1) + \mathrm{struct}(1) = n_{\mathrm{base}}$, exact on 6 shells.
6. **Phantom saturation C2**: for $|B| \geq 3$, the phantom absorbs all cross-terms.
7. **PM–PT corollary**: $\mathrm{AT} = 1$ for all $p \geq 41$, proving that structural complexity is exhausted in the first primorial cycle.

The constant $N_{\mathrm{spatial}} = 3$ appearing in the AT formula is the same quantity that determines spatial dimensionality via the active prime criterion [3]: it is the obstruction constant for structural detection through arithmetic noise.

Convergence extensions. The arithmetic–geometric race reveals five alternative routes to T4, all consistent. The mod-3 transition graph K_3 is an optimal Ramanujan expander ($\lambda_2(L) = 3$, Cheeger constant $h = 1$), and its Ihara zeta function places all non-trivial zeros on the critical circle $|u| = 1/\sqrt{q}$. The dihedral group $D_4 = \langle T_3, J\rangle$ organises the spectral structure, with the intertwiner J exchanging the two eigensectors v_+ and v_-.

Coding-theoretic interpretation. The sieve is simultaneously a prime-generating algorithm, a CPTP decoherence channel, an error-correcting code with $d = n - 1$, and a spectral contraction with gap $1 - |\lambda_2|$. The code distance $d = n - 1$ is the arithmetic manifestation of informational persistence: the sieve produces a code whose redundancy grows without bound, ensuring that the prime pattern is recoverable from partial information. The CRT quantum code $[[3, 5.58, 1]]$ with Knill–Laflamme deviation $\Delta_{\mathrm{KL}} = 0.008$ shows that prediction and error correction are the same mathematical structure viewed from different angles.

Open question. The general position conjecture (Conjecture 4.8)—that no accidental linear relations arise among the CRT cross-terms of the observation matrix—is verified numerically for all 13 explored shells but not derived from arithmetic first principles. A proof would promote the dimension formula and activation spectrum from conditional to unconditional theorems.

Forward. The companion article [5] constructs the bridge from arithmetic to physics: Pontryagin duality on the CRT decomposition, Buchstab contraction, and Osterwalder–Schrader reconstruction provide a formal pathway from the sieve structures developed here to the product coupling $\alpha_{\text{sieve}} = \prod_{p\in\{3,5,7\}} \sin^2\theta_p$ and its physical interpretation.

References

[1] G. H. Hardy, J. E. Littlewood, and G. Pólya. *Inequalities.* Cambridge University Press, 2 edition, 1952. First edition 1934.

[2] Yan Senez. The theory of persistence I: The sieve as a dynamical system. Companion article (M1), 2026.

[3] Yan Senez. The theory of persistence II: Information geometry and holonomy. Companion article (M2), 2026.

[4] Yan Senez. The theory of persistence III: Convergence and the fixed point. Companion article (M3), 2026.

[5] Yan Senez. The theory of persistence V: The bridge from arithmetic to physics. Companion article (M5), 2026.

[6] Yan Senez. The theory of persistence: Chemistry and molecular observables. Companion article (chemistry), 2026.

[7] Yan Senez. The theory of persistence: Physical observables from prime gap arithmetic. Companion article (physics), 2026.